

```
In [1]: !pip3 install scikit-learn==1.3.2
!pip install imbalanced-learn==0.11.0
import pandas as pd
import numpy as np

%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.impute import SimpleImputer

from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import (
    AdaBoostClassifier,
    GradientBoostingClassifier,
    RandomForestClassifier,
    BaggingClassifier,
)
from xgboost import XGBClassifier

from sklearn import metrics
from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.metrics import (
    f1_score,
    accuracy_score,
    recall_score,
    precision_score,
    confusion_matrix,
    roc_auc_score,
    ConfusionMatrixDisplay,
)

from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder

from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer

pd.set_option("display.max_columns", None)

pd.set_option("display.float_format", lambda x: "%.3f" % x)

import warnings

warnings.filterwarnings("ignore")
```

Requirement already satisfied: scikit-learn==1.3.2 in c:\users\conne\anaconda3\lib\site-packages (1.3.2)
 Requirement already satisfied: numpy<2.0,>=1.17.3 in c:\users\conne\anaconda3\lib\site-packages (from scikit-learn==1.3.2) (1.24.3)
 Requirement already satisfied: scipy>=1.5.0 in c:\users\conne\anaconda3\lib\site-packages (from scikit-learn==1.3.2) (1.10.1)
 Requirement already satisfied: joblib>=1.1.1 in c:\users\conne\anaconda3\lib\site-packages (from scikit-learn==1.3.2) (1.2.0)
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\conne\anaconda3\lib\site-packages (from scikit-learn==1.3.2) (2.2.0)
 Requirement already satisfied: imbalanced-learn==0.11.0 in c:\users\conne\anaconda3\lib\site-packages (0.11.0)
 Requirement already satisfied: numpy>=1.17.3 in c:\users\conne\anaconda3\lib\site-packages (from imbalanced-learn==0.11.0) (1.24.3)
 Requirement already satisfied: scipy>=1.5.0 in c:\users\conne\anaconda3\lib\site-packages (from imbalanced-learn==0.11.0) (1.10.1)
 Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\conne\anaconda3\lib\site-packages (from imbalanced-learn==0.11.0) (1.3.2)
 Requirement already satisfied: joblib>=1.1.1 in c:\users\conne\anaconda3\lib\site-packages (from imbalanced-learn==0.11.0) (1.2.0)
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\conne\anaconda3\lib\site-packages (from imbalanced-learn==0.11.0) (2.2.0)

Here im importing all packages

```
In [2]: df = pd.read_csv('/Users/conne/Downloads/Train.csv.csv')
df_test = pd.read_csv('/Users/conne/Downloads/Test.csv.csv')
```

Here im loading both sets of data

```
In [3]: data = df.copy()
```

```
In [4]: data_test = df_test.copy()
```

```
In [5]: data.shape
```

```
Out[5]: (20000, 41)
```

We can see the first set has 20000 rows and 41 columns

```
In [6]: data_test.shape
```

```
Out[6]: (5000, 41)
```

We can see the first set has 5000 rows and 41 columns

```
In [7]: data.head()
```

Out[7]:

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13
0	-4.465	-4.679	3.102	0.506	-0.221	-2.033	-2.911	0.051	-1.522	3.762	-5.715	0.736	0.981
1	3.366	3.653	0.910	-1.368	0.332	2.359	0.733	-4.332	0.566	-0.101	1.914	-0.951	-1.255
2	-3.832	-5.824	0.634	-2.419	-1.774	1.017	-2.099	-3.173	-2.082	5.393	-0.771	1.107	1.144
3	1.618	1.888	7.046	-1.147	0.083	-1.530	0.207	-2.494	0.345	2.119	-3.053	0.460	2.705
4	-0.111	3.872	-3.758	-2.983	3.793	0.545	0.205	4.849	-1.855	-6.220	1.998	4.724	0.709

In [8]: data.tail()

Out[8]:

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13
19995	-2.071	-1.088	-0.796	-3.012	-2.288	2.807	0.481	0.105	-0.587	-2.899	8.868	1.717	1.3
19996	2.890	2.483	5.644	0.937	-1.381	0.412	-1.593	-5.762	2.150	0.272	-2.095	-1.526	0.0
19997	-3.897	-3.942	-0.351	-2.417	1.108	-1.528	-3.520	2.055	-0.234	-0.358	-3.782	2.180	6.1
19998	-3.187	-10.052	5.696	-4.370	-5.355	-1.873	-3.947	0.679	-2.389	5.457	1.583	3.571	9.2
19999	-2.687	1.961	6.137	2.600	2.657	-4.291	-2.344	0.974	-1.027	0.497	-9.589	3.177	1.0

In [9]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 41 columns):
#   Column      Non-Null Count  Dtype
---  -
0   V1          19982 non-null  float64
1   V2          19982 non-null  float64
2   V3          20000 non-null  float64
3   V4          20000 non-null  float64
4   V5          20000 non-null  float64
5   V6          20000 non-null  float64
6   V7          20000 non-null  float64
7   V8          20000 non-null  float64
8   V9          20000 non-null  float64
9   V10         20000 non-null  float64
10  V11         20000 non-null  float64
11  V12         20000 non-null  float64
12  V13         20000 non-null  float64
13  V14         20000 non-null  float64
14  V15         20000 non-null  float64
15  V16         20000 non-null  float64
16  V17         20000 non-null  float64
17  V18         20000 non-null  float64
18  V19         20000 non-null  float64
19  V20         20000 non-null  float64
20  V21         20000 non-null  float64
21  V22         20000 non-null  float64
22  V23         20000 non-null  float64
23  V24         20000 non-null  float64
24  V25         20000 non-null  float64
25  V26         20000 non-null  float64
26  V27         20000 non-null  float64
27  V28         20000 non-null  float64
28  V29         20000 non-null  float64
29  V30         20000 non-null  float64
30  V31         20000 non-null  float64
31  V32         20000 non-null  float64
32  V33         20000 non-null  float64
33  V34         20000 non-null  float64
34  V35         20000 non-null  float64
35  V36         20000 non-null  float64
36  V37         20000 non-null  float64
37  V38         20000 non-null  float64
38  V39         20000 non-null  float64
39  V40         20000 non-null  float64
40  Target      20000 non-null  int64
dtypes: float64(40), int64(1)
memory usage: 6.3 MB
```

```
In [10]: data.duplicated().sum()
```

```
Out[10]: 0
```

```
In [11]: data.isnull().sum()
```

```
Out[11]: V1      18
          V2      18
          V3       0
          V4       0
          V5       0
          V6       0
          V7       0
          V8       0
          V9       0
          V10      0
          V11      0
          V12      0
          V13      0
          V14      0
          V15      0
          V16      0
          V17      0
          V18      0
          V19      0
          V20      0
          V21      0
          V22      0
          V23      0
          V24      0
          V25      0
          V26      0
          V27      0
          V28      0
          V29      0
          V30      0
          V31      0
          V32      0
          V33      0
          V34      0
          V35      0
          V36      0
          V37      0
          V38      0
          V39      0
          V40      0
          Target    0
          dtype: int64
```

we have a couple null values that must be fixed

```
In [12]: data_test.isnull().sum()
```

```
Out[12]: V1      5
          V2      6
          V3      0
          V4      0
          V5      0
          V6      0
          V7      0
          V8      0
          V9      0
          V10     0
          V11     0
          V12     0
          V13     0
          V14     0
          V15     0
          V16     0
          V17     0
          V18     0
          V19     0
          V20     0
          V21     0
          V22     0
          V23     0
          V24     0
          V25     0
          V26     0
          V27     0
          V28     0
          V29     0
          V30     0
          V31     0
          V32     0
          V33     0
          V34     0
          V35     0
          V36     0
          V37     0
          V38     0
          V39     0
          V40     0
          Target   0
          dtype: int64
```

we have a couple null values that must be fixed

```
In [13]: data.describe().T
```

Out[13]:

	count	mean	std	min	25%	50%	75%	max
V1	19982.000	-0.272	3.442	-11.876	-2.737	-0.748	1.840	15.493
V2	19982.000	0.440	3.151	-12.320	-1.641	0.472	2.544	13.089
V3	20000.000	2.485	3.389	-10.708	0.207	2.256	4.566	17.091
V4	20000.000	-0.083	3.432	-15.082	-2.348	-0.135	2.131	13.236
V5	20000.000	-0.054	2.105	-8.603	-1.536	-0.102	1.340	8.134
V6	20000.000	-0.995	2.041	-10.227	-2.347	-1.001	0.380	6.976
V7	20000.000	-0.879	1.762	-7.950	-2.031	-0.917	0.224	8.006
V8	20000.000	-0.548	3.296	-15.658	-2.643	-0.389	1.723	11.679
V9	20000.000	-0.017	2.161	-8.596	-1.495	-0.068	1.409	8.138
V10	20000.000	-0.013	2.193	-9.854	-1.411	0.101	1.477	8.108
V11	20000.000	-1.895	3.124	-14.832	-3.922	-1.921	0.119	11.826
V12	20000.000	1.605	2.930	-12.948	-0.397	1.508	3.571	15.081
V13	20000.000	1.580	2.875	-13.228	-0.224	1.637	3.460	15.420
V14	20000.000	-0.951	1.790	-7.739	-2.171	-0.957	0.271	5.671
V15	20000.000	-2.415	3.355	-16.417	-4.415	-2.383	-0.359	12.246
V16	20000.000	-2.925	4.222	-20.374	-5.634	-2.683	-0.095	13.583
V17	20000.000	-0.134	3.345	-14.091	-2.216	-0.015	2.069	16.756
V18	20000.000	1.189	2.592	-11.644	-0.404	0.883	2.572	13.180
V19	20000.000	1.182	3.397	-13.492	-1.050	1.279	3.493	13.238
V20	20000.000	0.024	3.669	-13.923	-2.433	0.033	2.512	16.052
V21	20000.000	-3.611	3.568	-17.956	-5.930	-3.533	-1.266	13.840
V22	20000.000	0.952	1.652	-10.122	-0.118	0.975	2.026	7.410
V23	20000.000	-0.366	4.032	-14.866	-3.099	-0.262	2.452	14.459
V24	20000.000	1.134	3.912	-16.387	-1.468	0.969	3.546	17.163
V25	20000.000	-0.002	2.017	-8.228	-1.365	0.025	1.397	8.223
V26	20000.000	1.874	3.435	-11.834	-0.338	1.951	4.130	16.836
V27	20000.000	-0.612	4.369	-14.905	-3.652	-0.885	2.189	17.560
V28	20000.000	-0.883	1.918	-9.269	-2.171	-0.891	0.376	6.528
V29	20000.000	-0.986	2.684	-12.579	-2.787	-1.176	0.630	10.722
V30	20000.000	-0.016	3.005	-14.796	-1.867	0.184	2.036	12.506
V31	20000.000	0.487	3.461	-13.723	-1.818	0.490	2.731	17.255
V32	20000.000	0.304	5.500	-19.877	-3.420	0.052	3.762	23.633
V33	20000.000	0.050	3.575	-16.898	-2.243	-0.066	2.255	16.692
V34	20000.000	-0.463	3.184	-17.985	-2.137	-0.255	1.437	14.358

	count	mean	std	min	25%	50%	75%	max
V35	20000.000	2.230	2.937	-15.350	0.336	2.099	4.064	15.291
V36	20000.000	1.515	3.801	-14.833	-0.944	1.567	3.984	19.330
V37	20000.000	0.011	1.788	-5.478	-1.256	-0.128	1.176	7.467
V38	20000.000	-0.344	3.948	-17.375	-2.988	-0.317	2.279	15.290
V39	20000.000	0.891	1.753	-6.439	-0.272	0.919	2.058	7.760
V40	20000.000	-0.876	3.012	-11.024	-2.940	-0.921	1.120	10.654
Target	20000.000	0.056	0.229	0.000	0.000	0.000	0.000	1.000

we can see the statistical summary

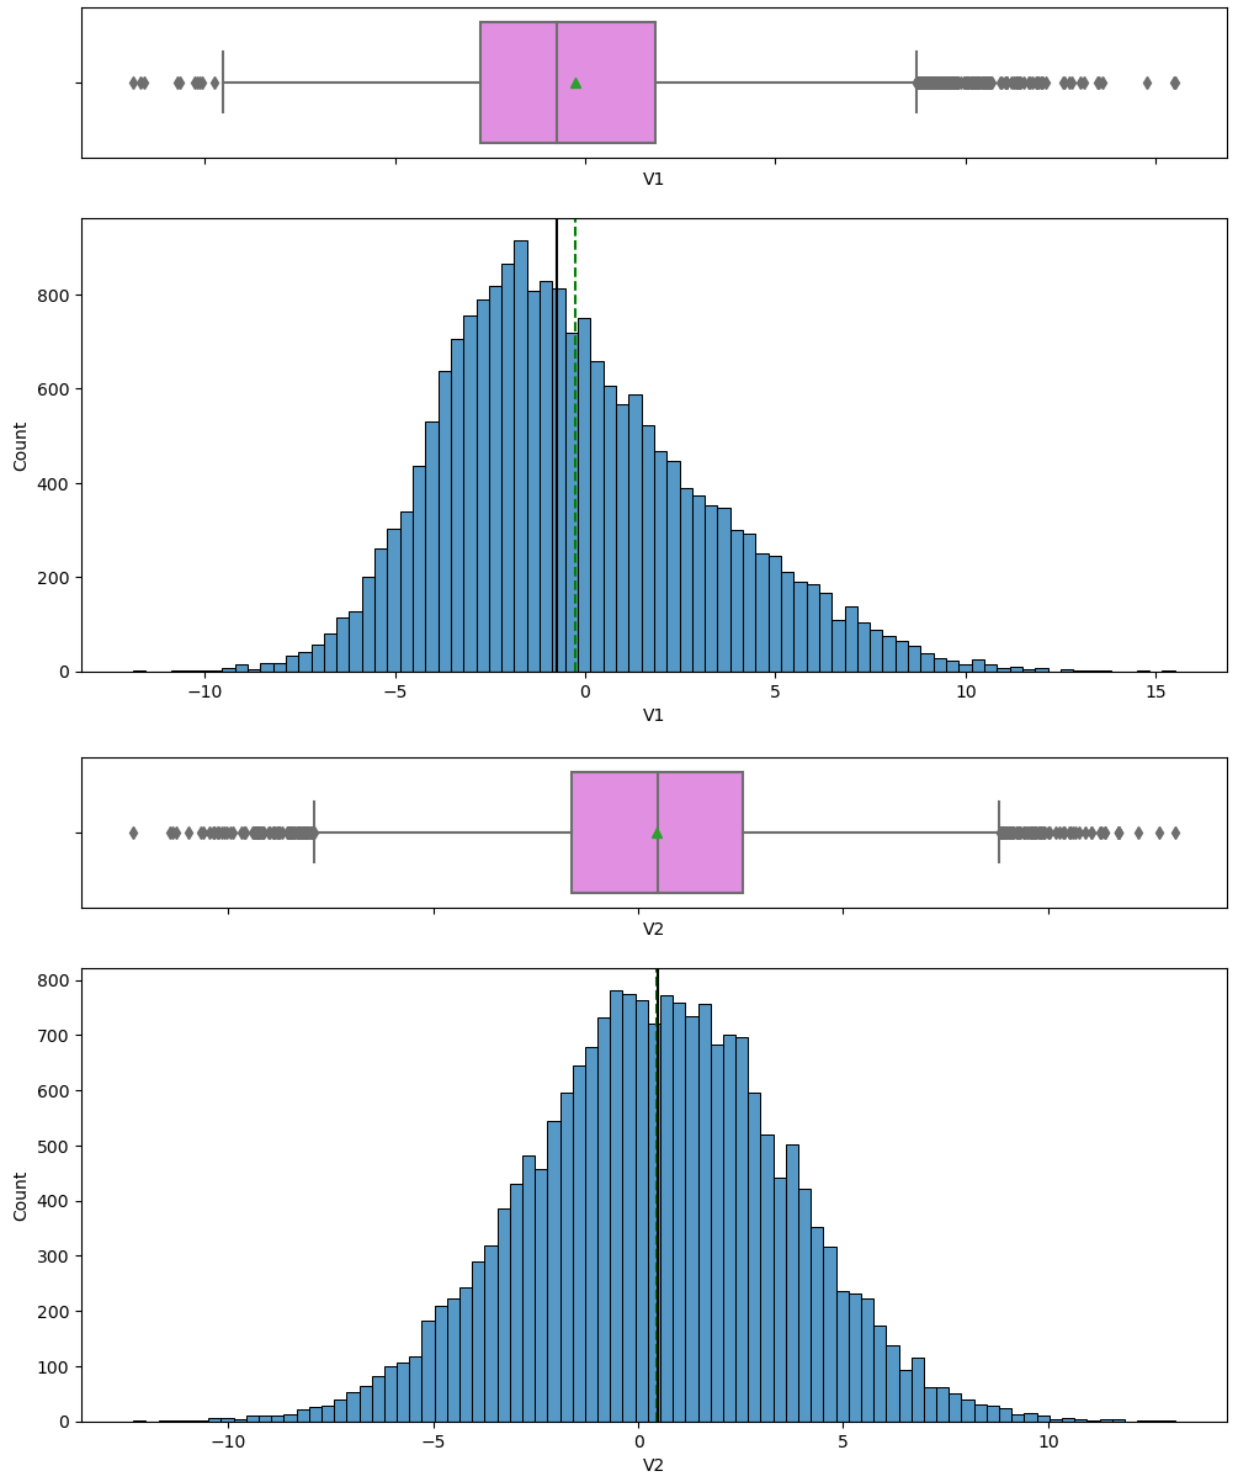
```
In [14]: def histogram_boxplot(data, feature, figsize=(12, 7), kde=False, bins=None):
        """
        Boxplot and histogram combined

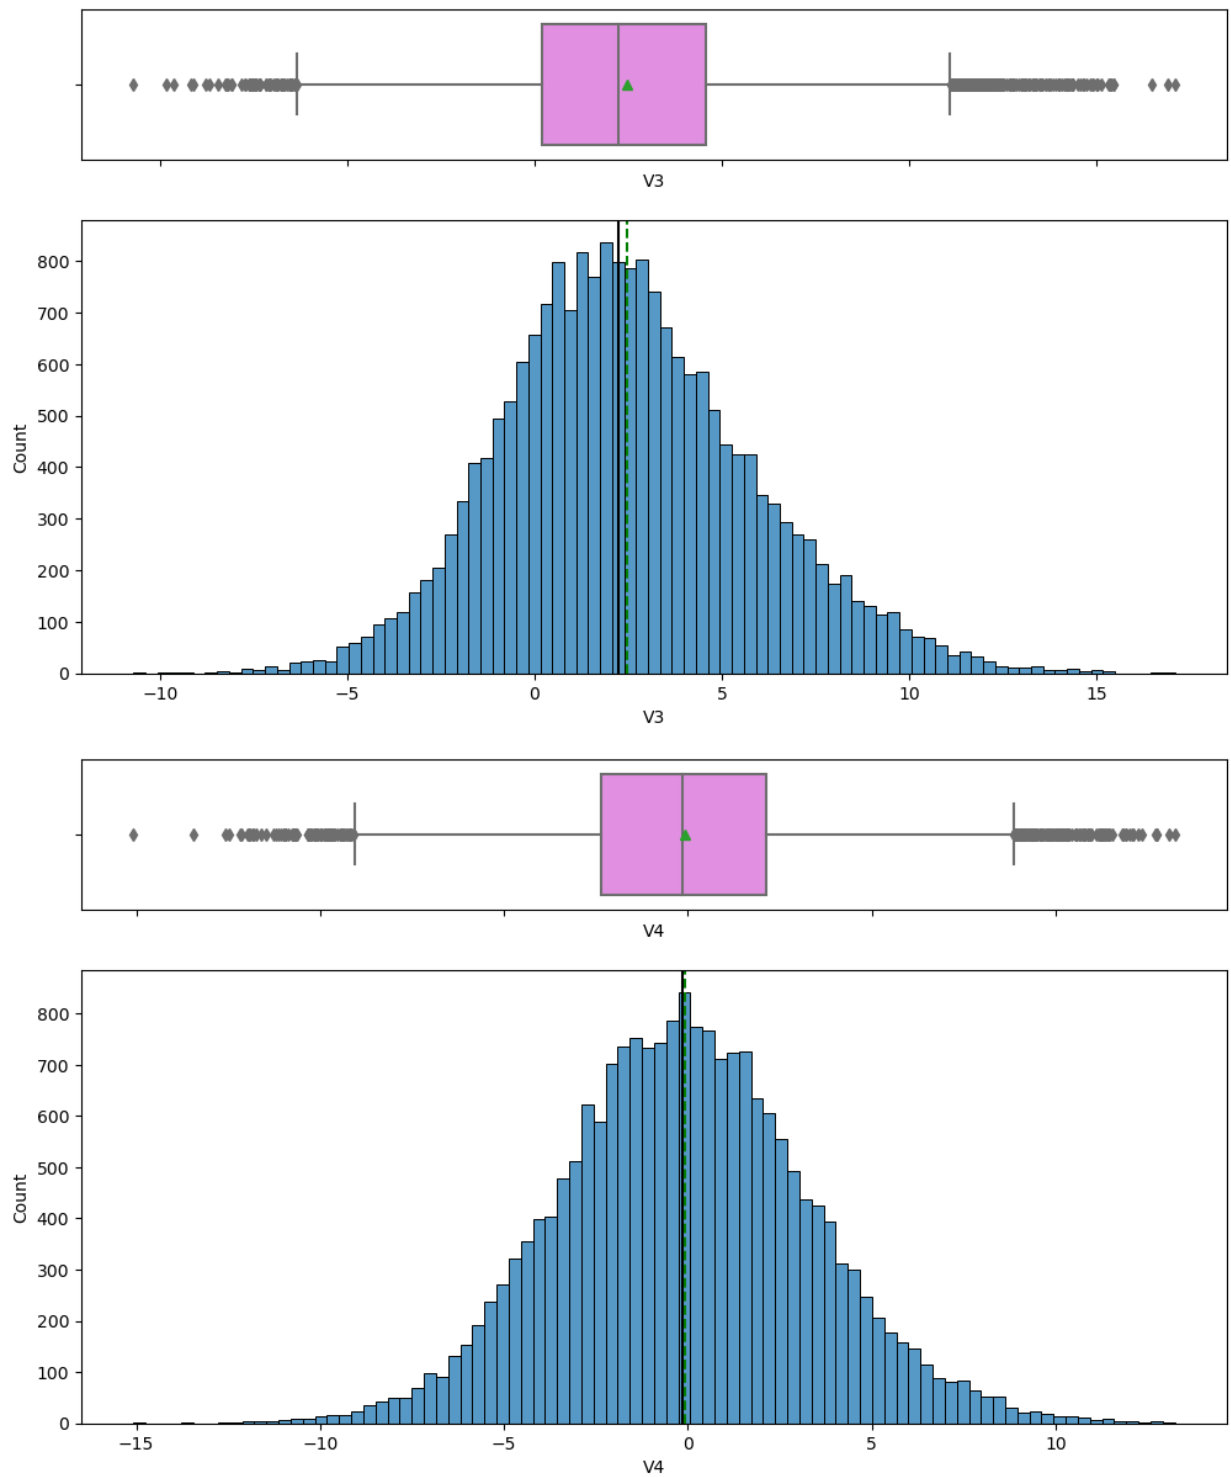
        data: dataframe
        feature: dataframe column
        figsize: size of figure (default (12,7))
        kde: whether to the show density curve (default False)
        bins: number of bins for histogram (default None)
        """

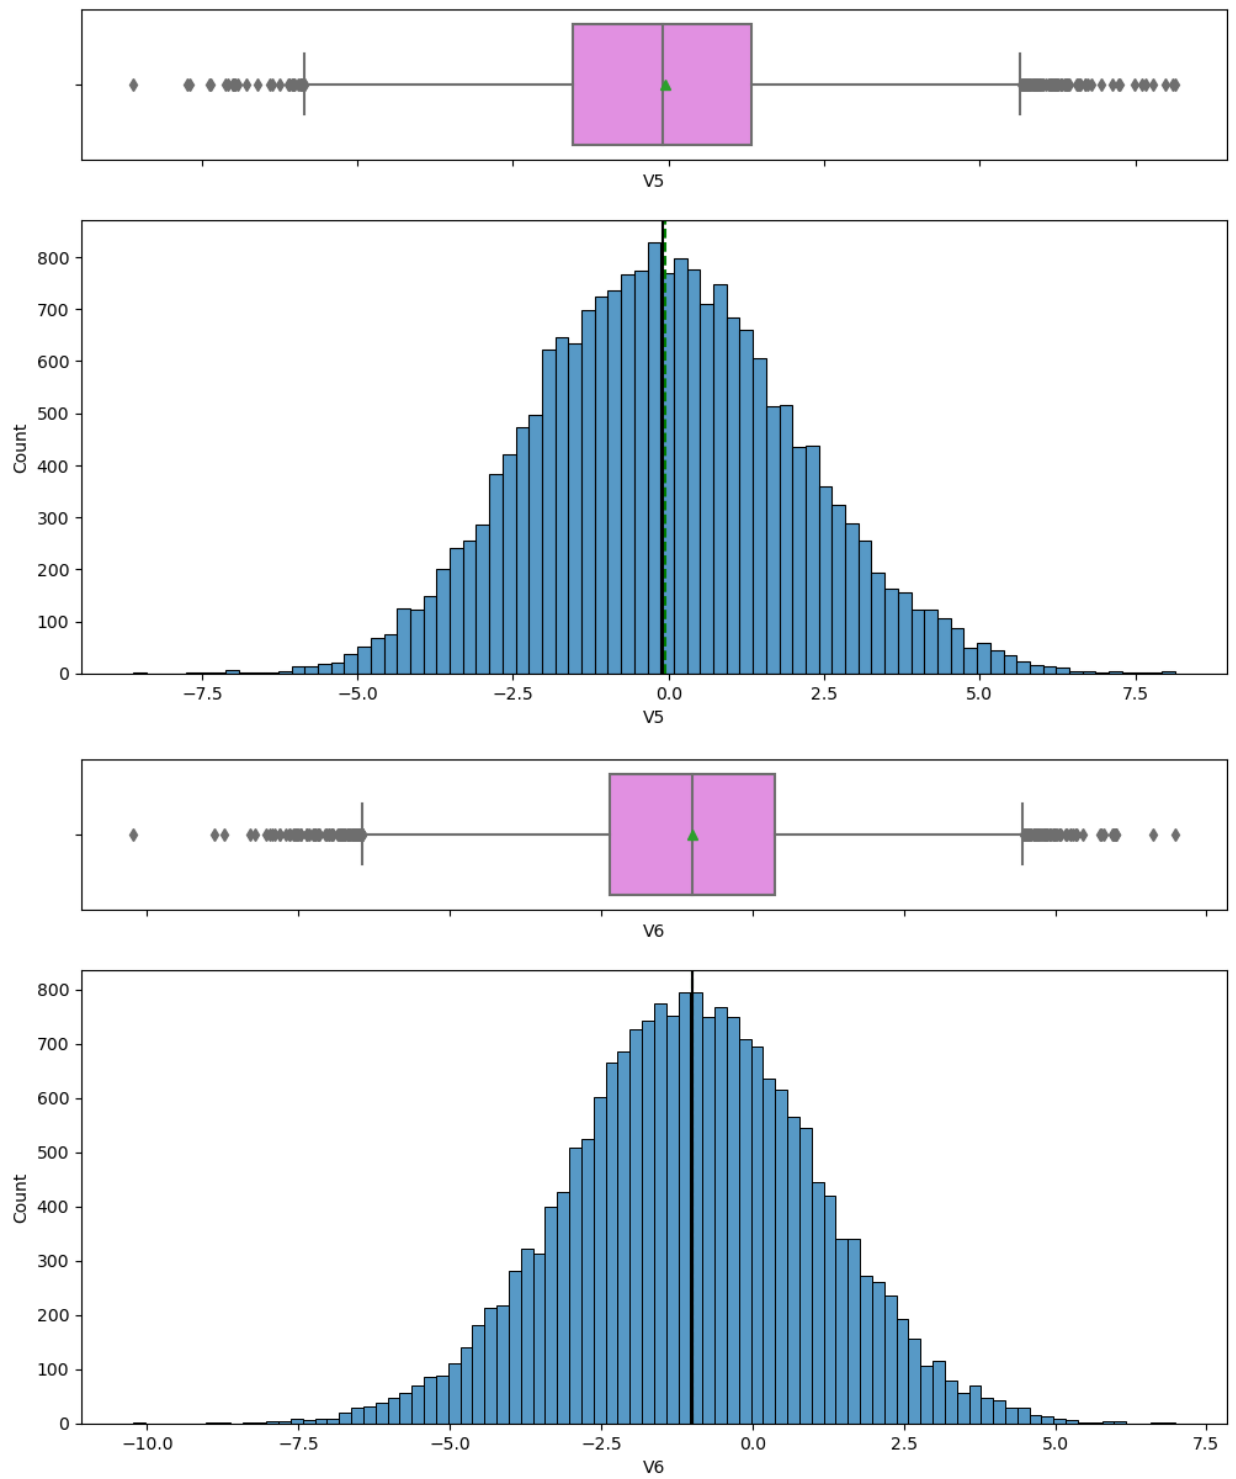
        f2, (ax_box2, ax_hist2) = plt.subplots(
            nrows=2, # Number of rows of the subplot grid= 2
            sharex=True, # x-axis will be shared among all subplots
            gridspec_kw={"height_ratios": (0.25, 0.75)},
            figsize=figsize,
        ) # creating the 2 subplots
        sns.boxplot(
            data=data, x=feature, ax=ax_box2, showmeans=True, color="violet"
        ) # boxplot will be created and a star will indicate the mean value of the column
        sns.histplot(
            data=data, x=feature, kde=kde, ax=ax_hist2, bins=bins, palette="winter"
        ) if bins else sns.histplot(
            data=data, x=feature, kde=kde, ax=ax_hist2
        ) # For histogram
        ax_hist2.axvline(
            data[feature].mean(), color="green", linestyle="--"
        ) # Add mean to the histogram
        ax_hist2.axvline(
            data[feature].median(), color="black", linestyle="-"
        ) # Add median to the histogram
```

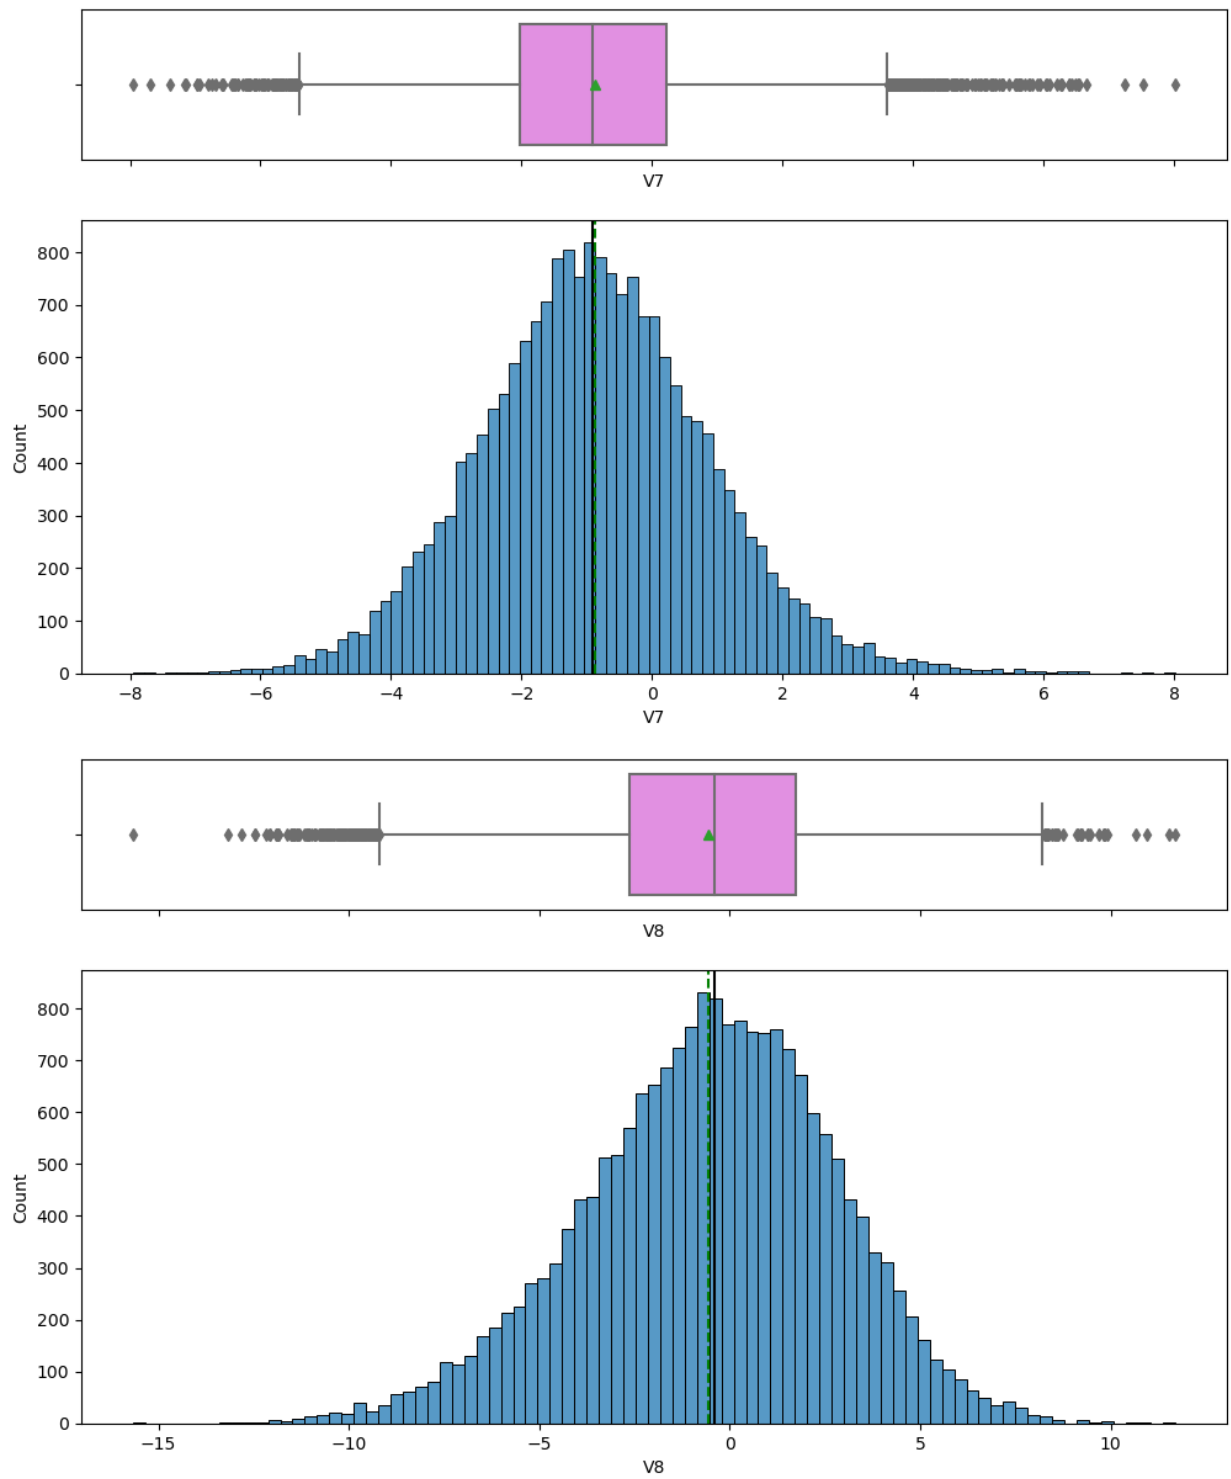
Here i will preform univariant analysis

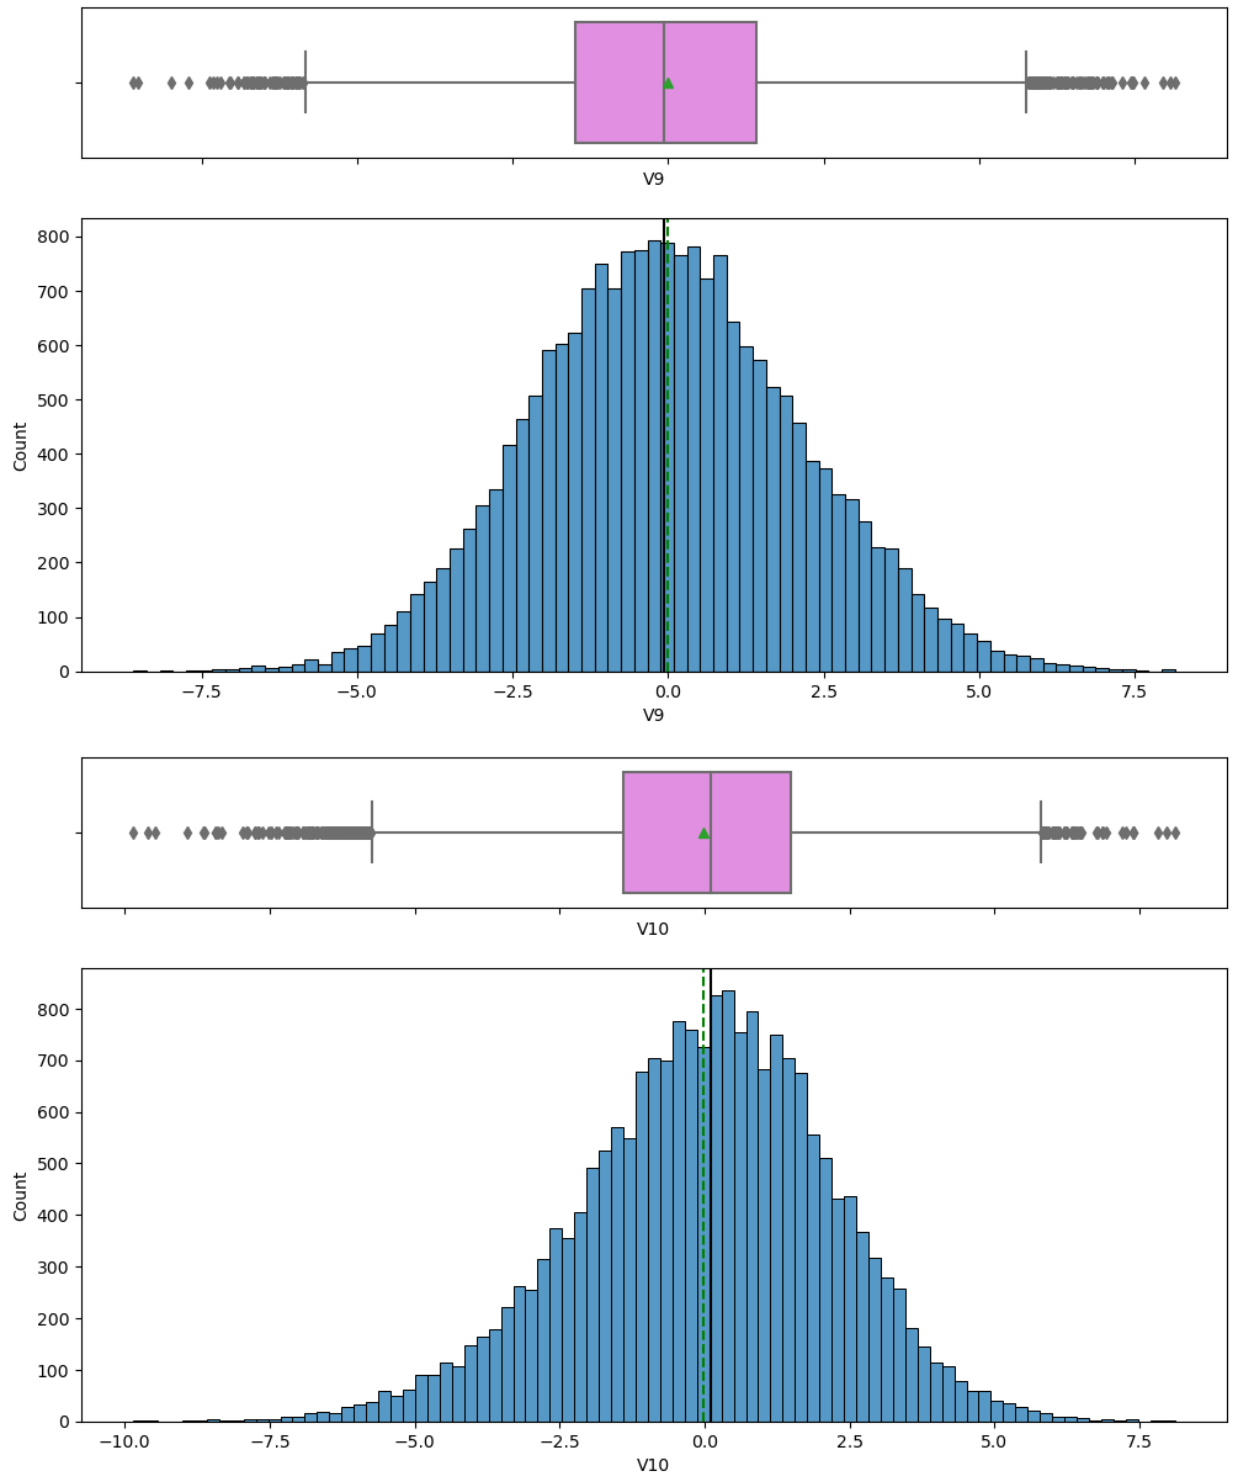
```
In [15]: for feature in df.columns :
        histogram_boxplot(df, feature, figsize=(12,7), kde=False, bins=None)
```

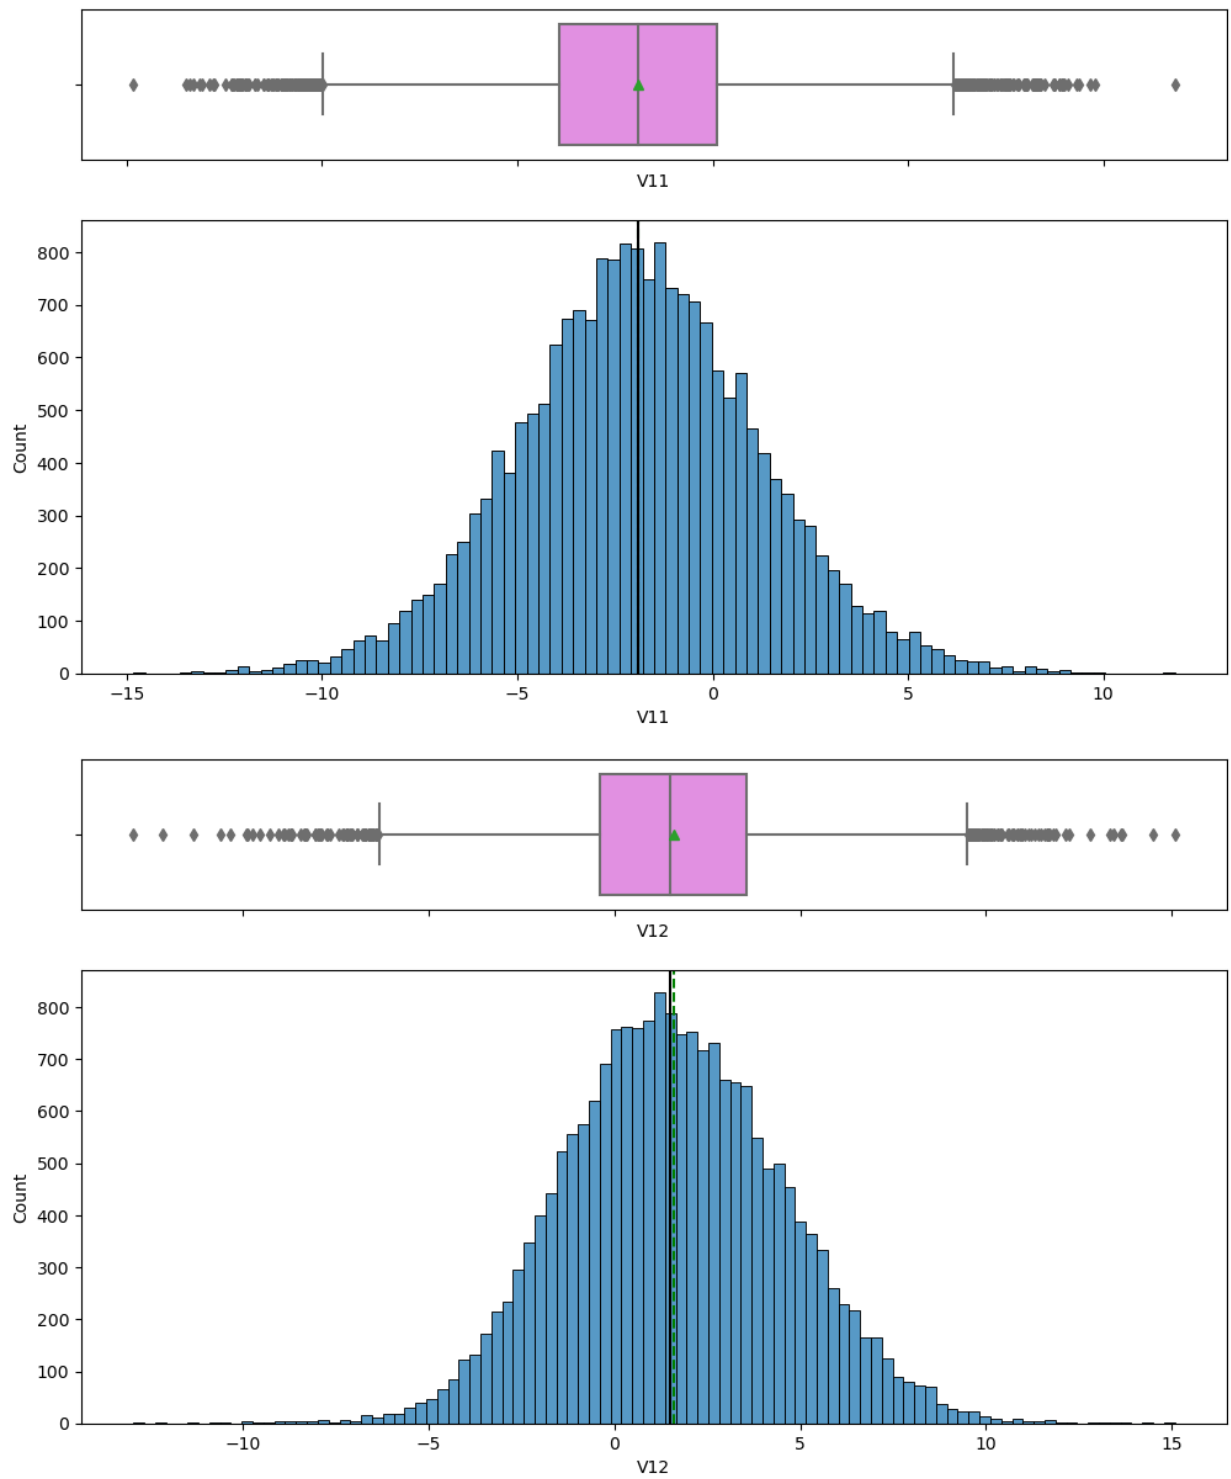



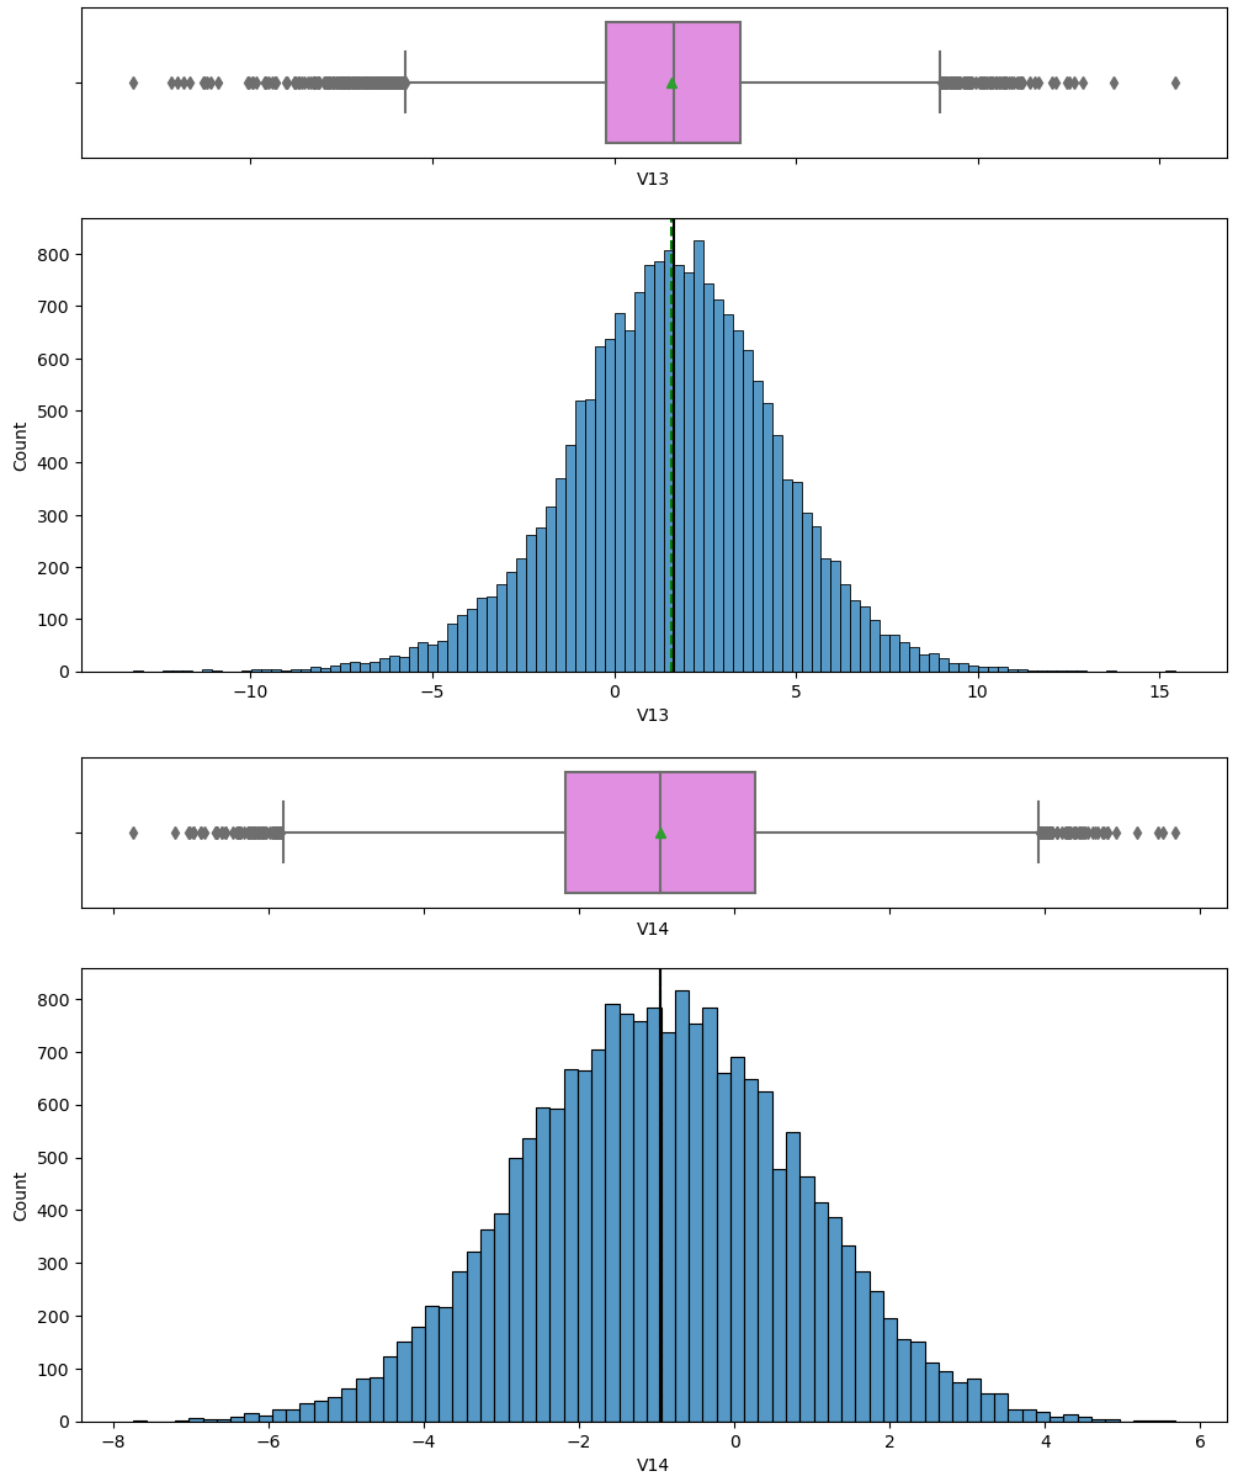


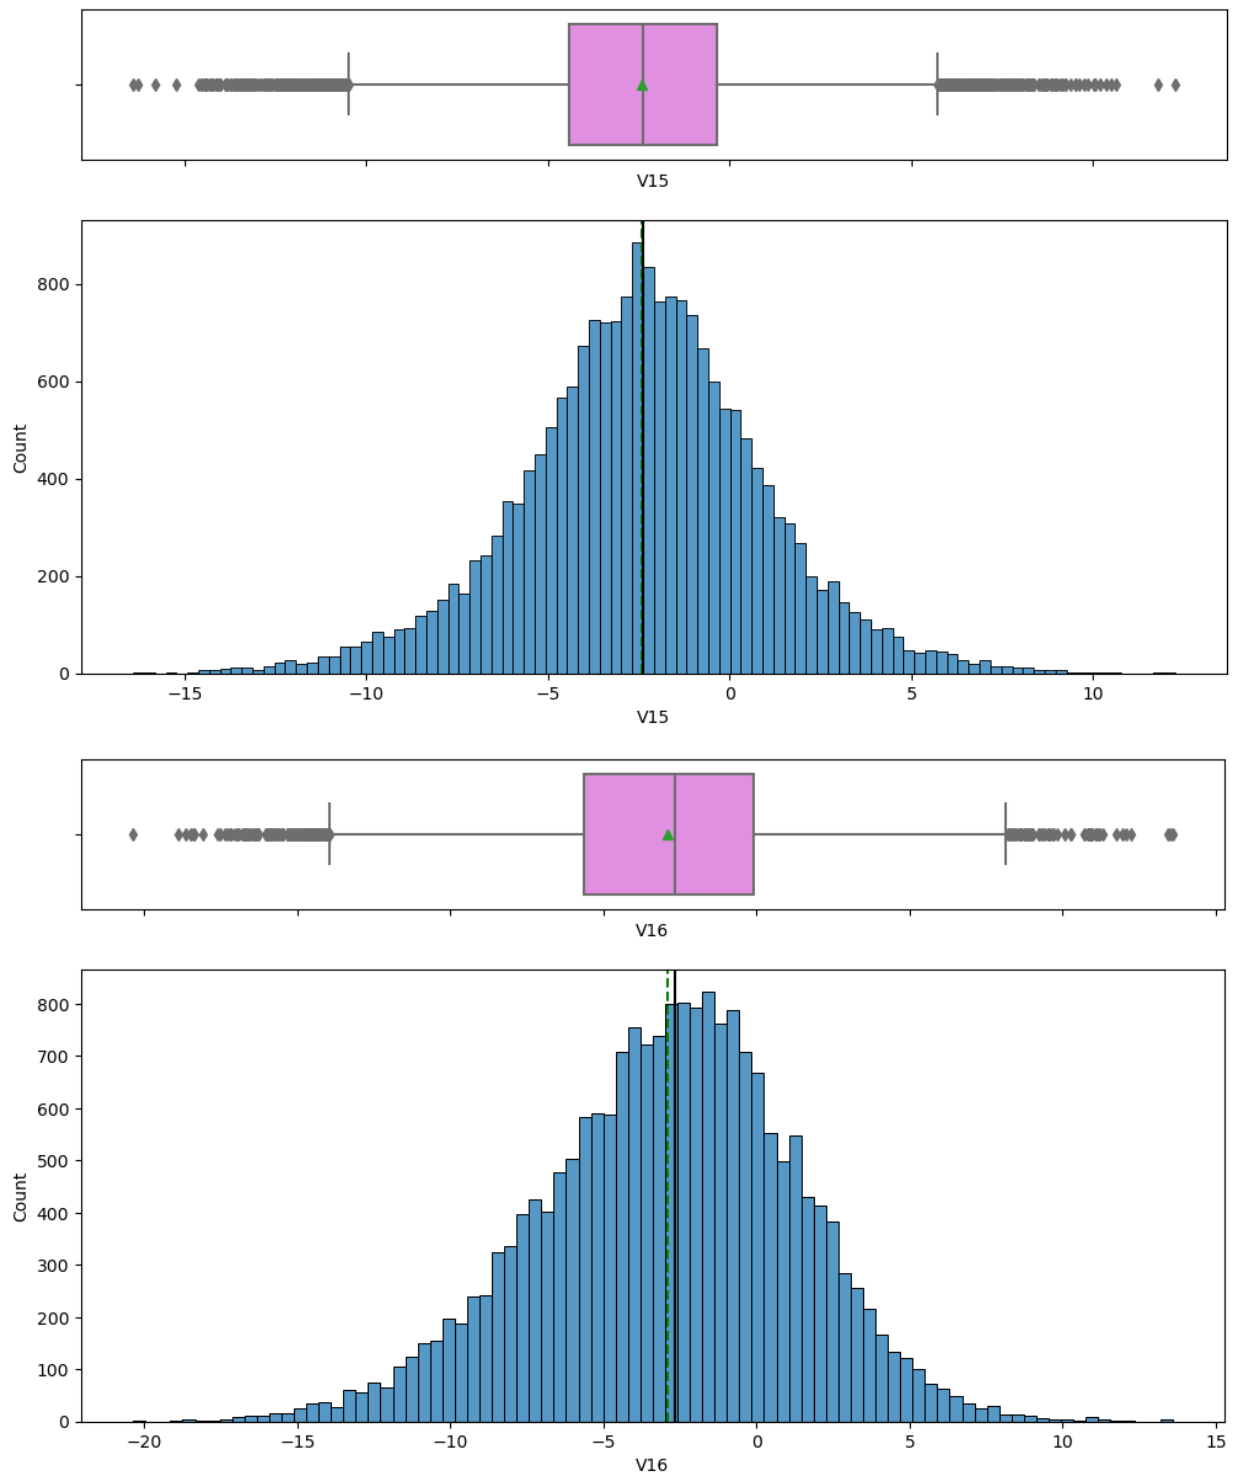


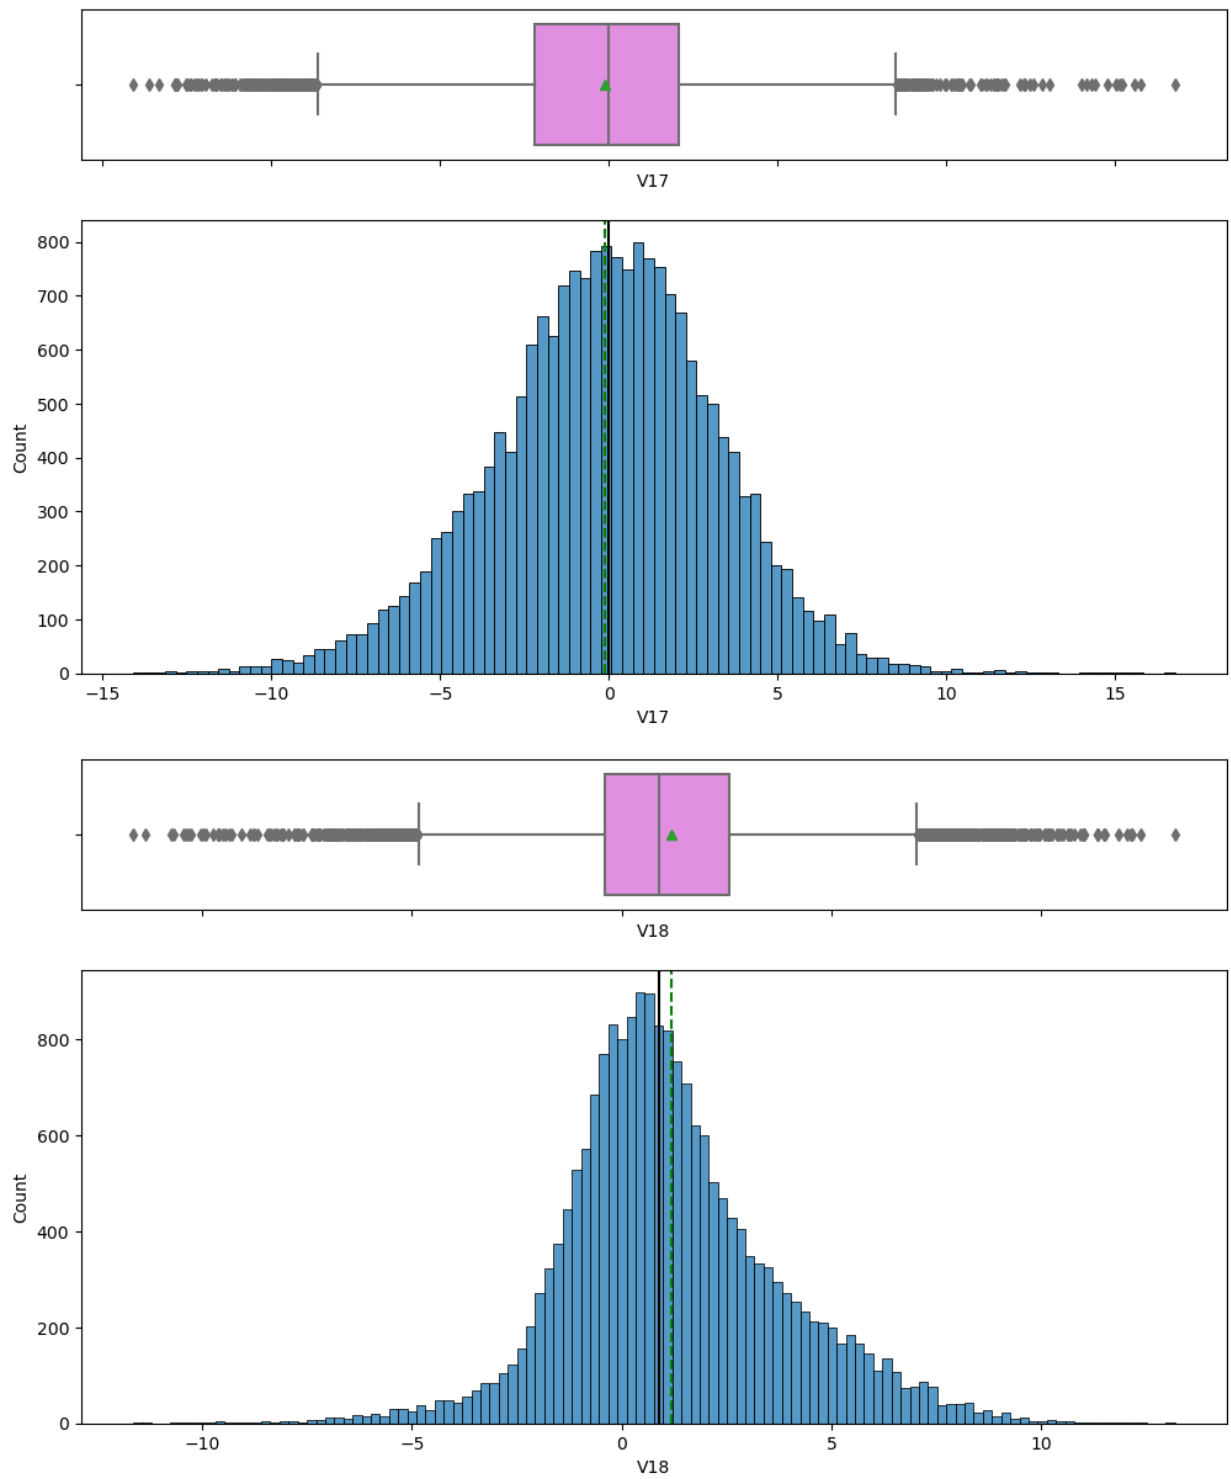


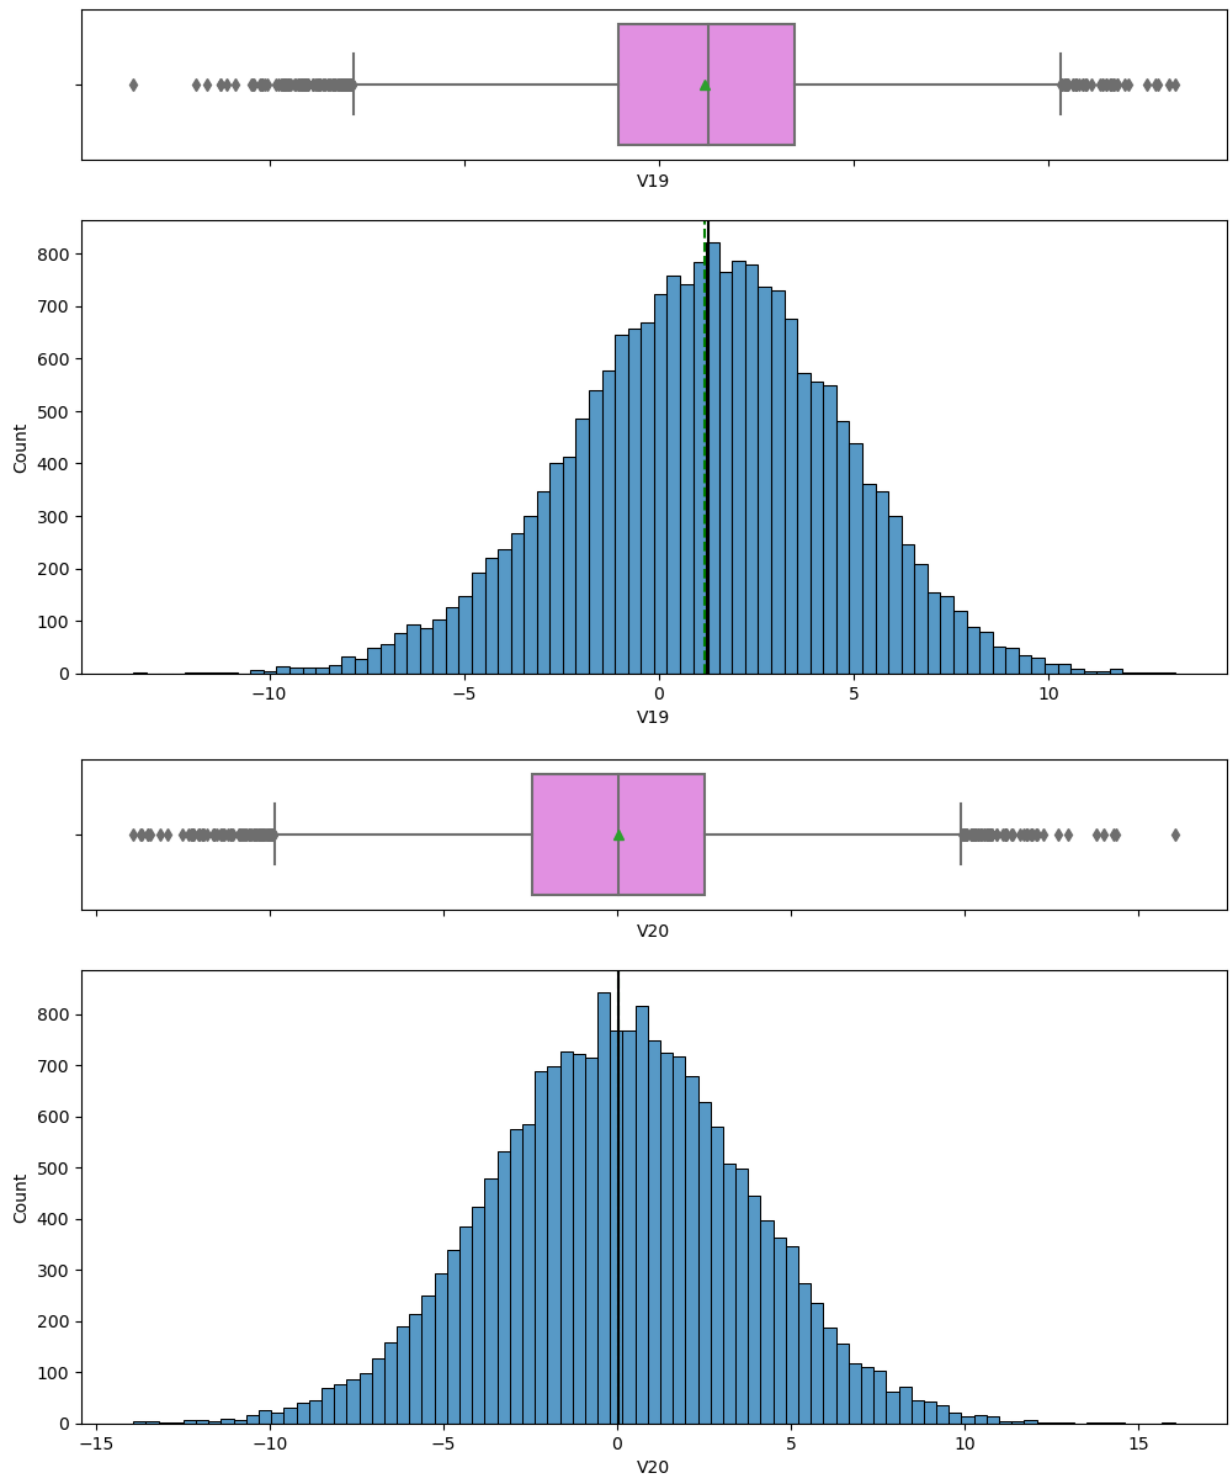


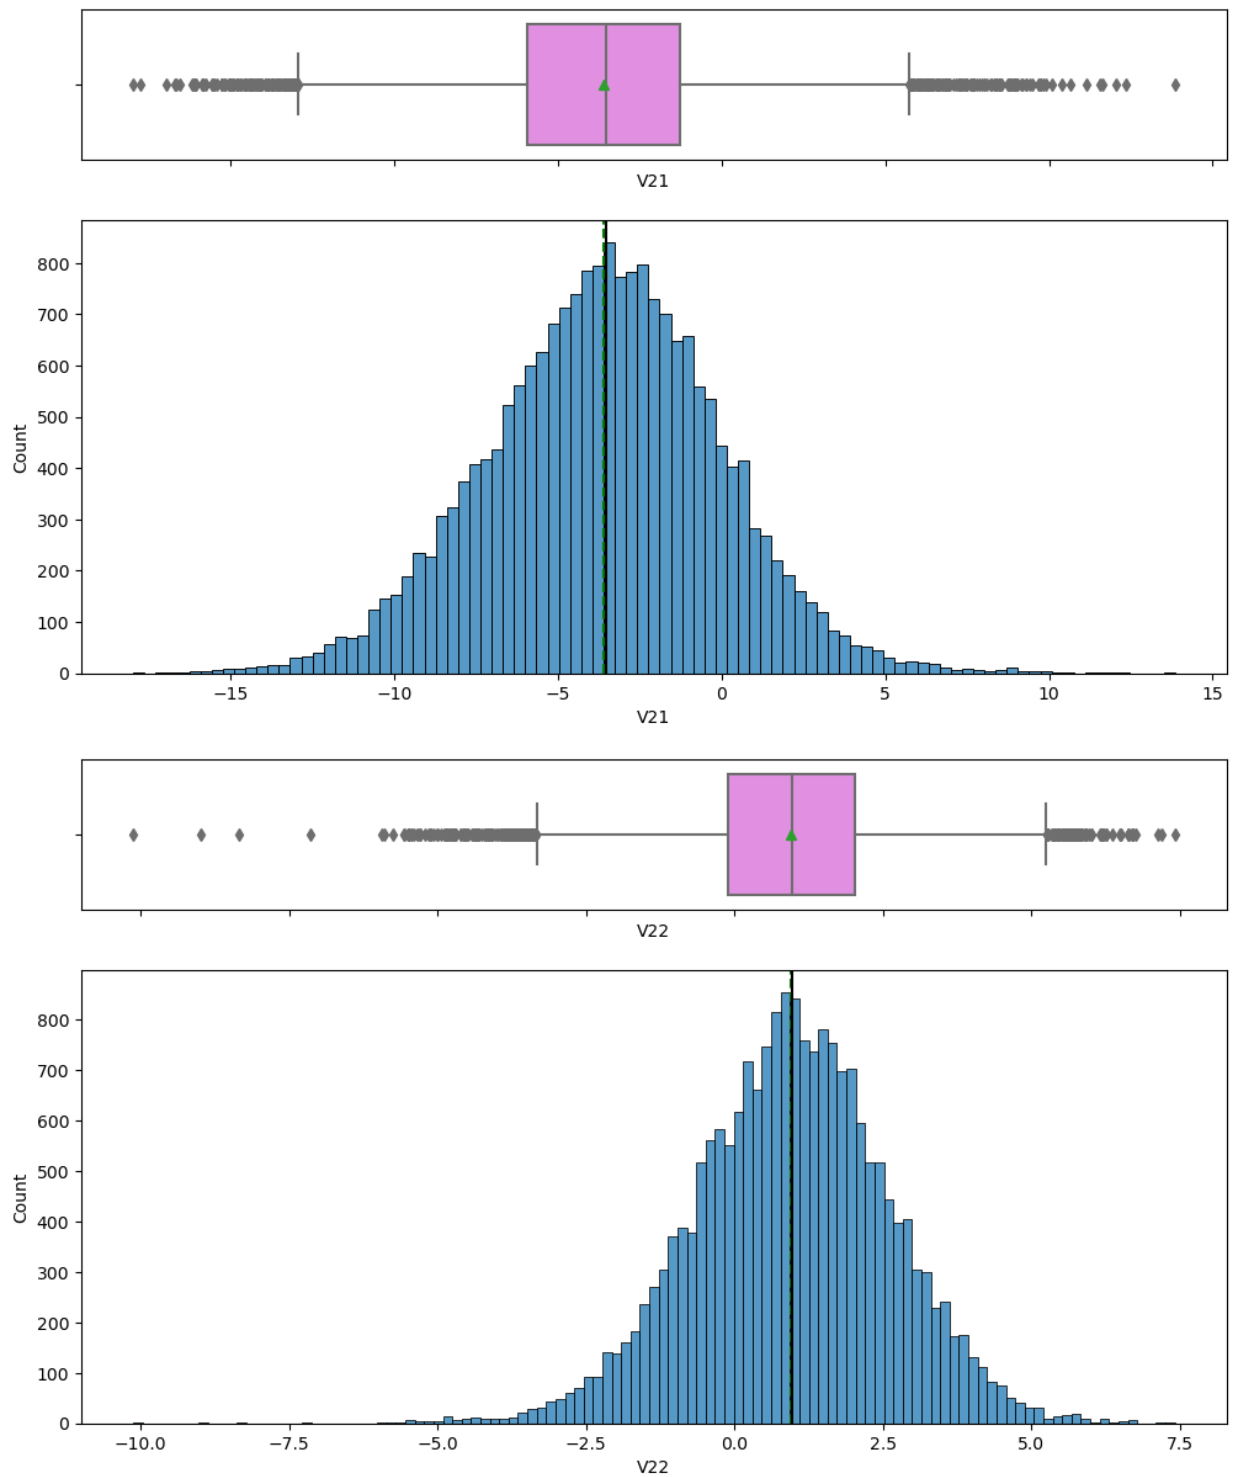


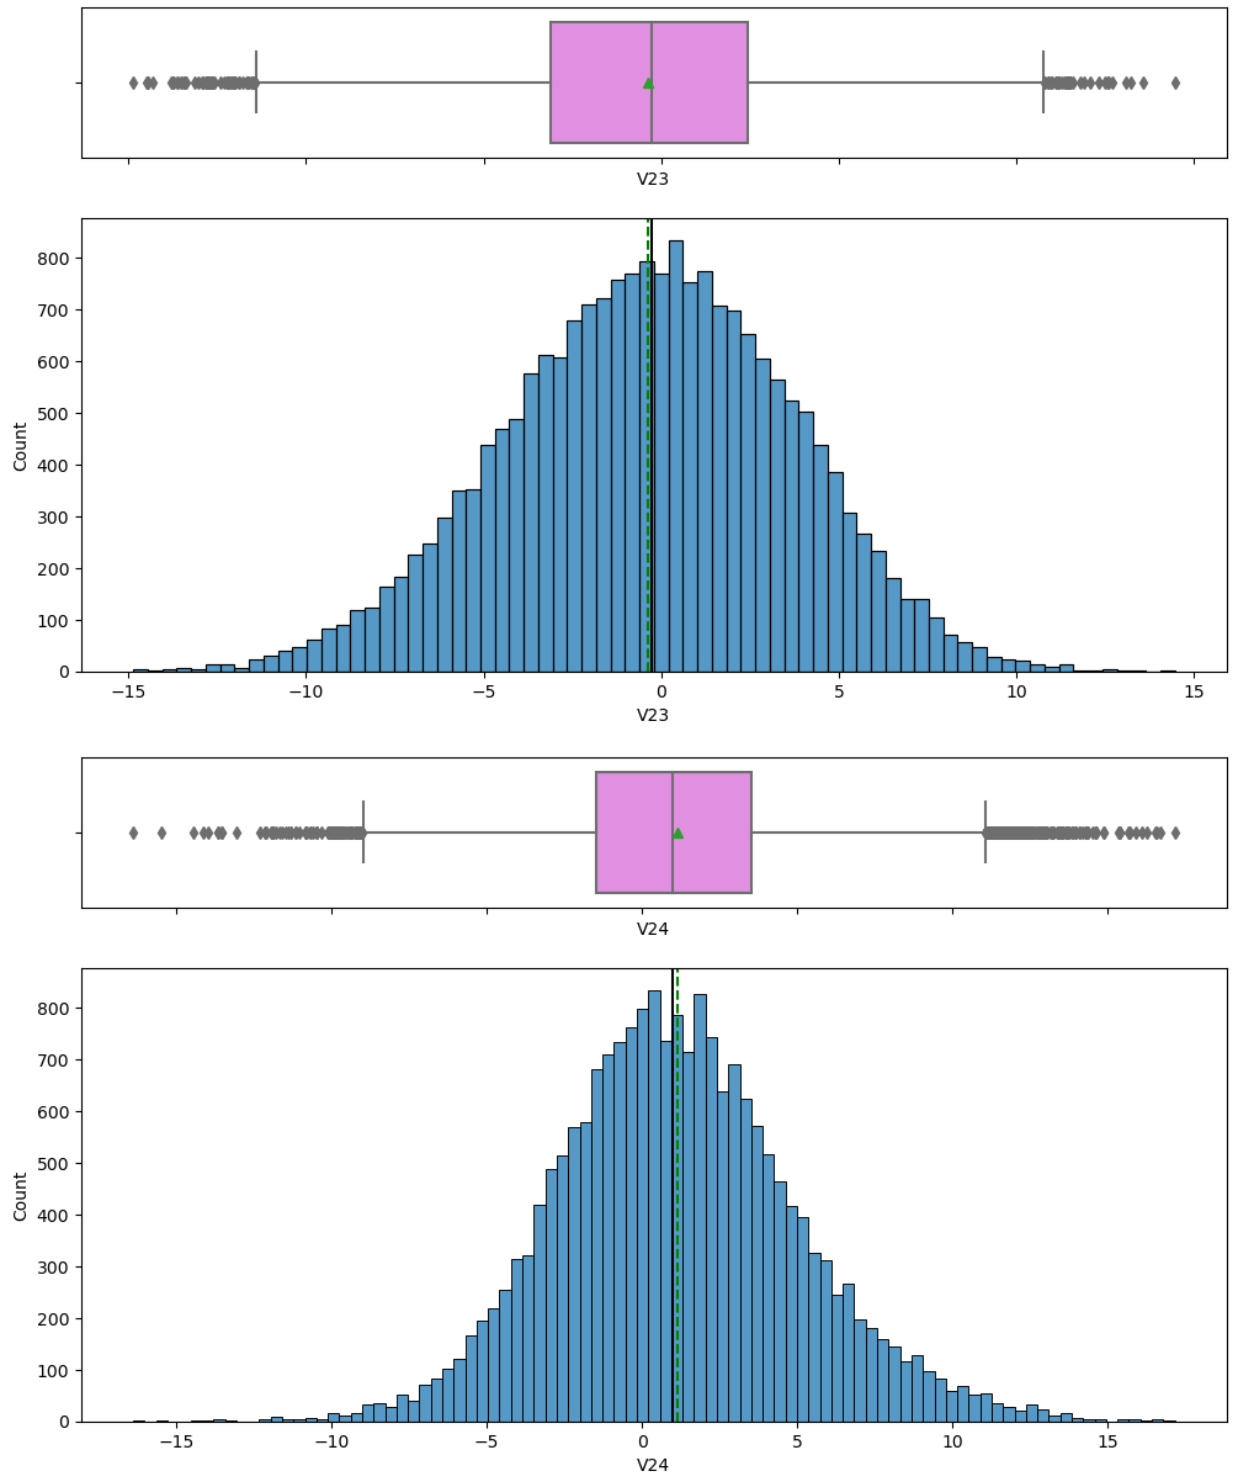


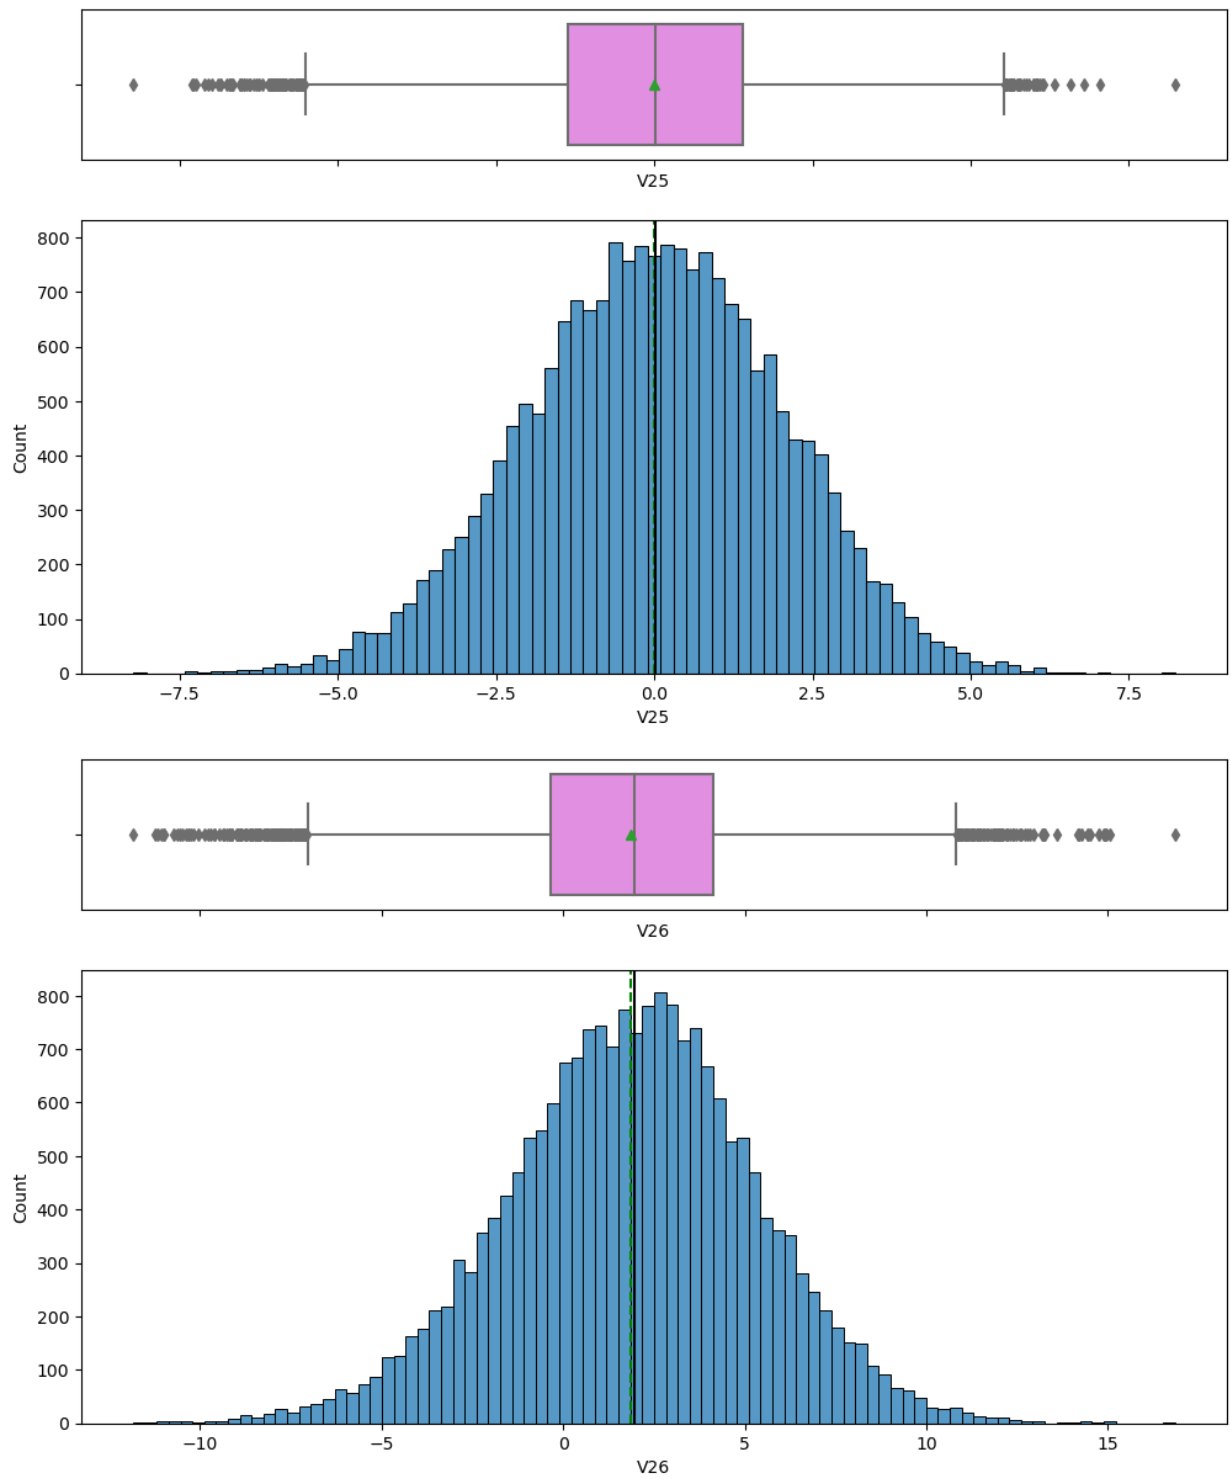


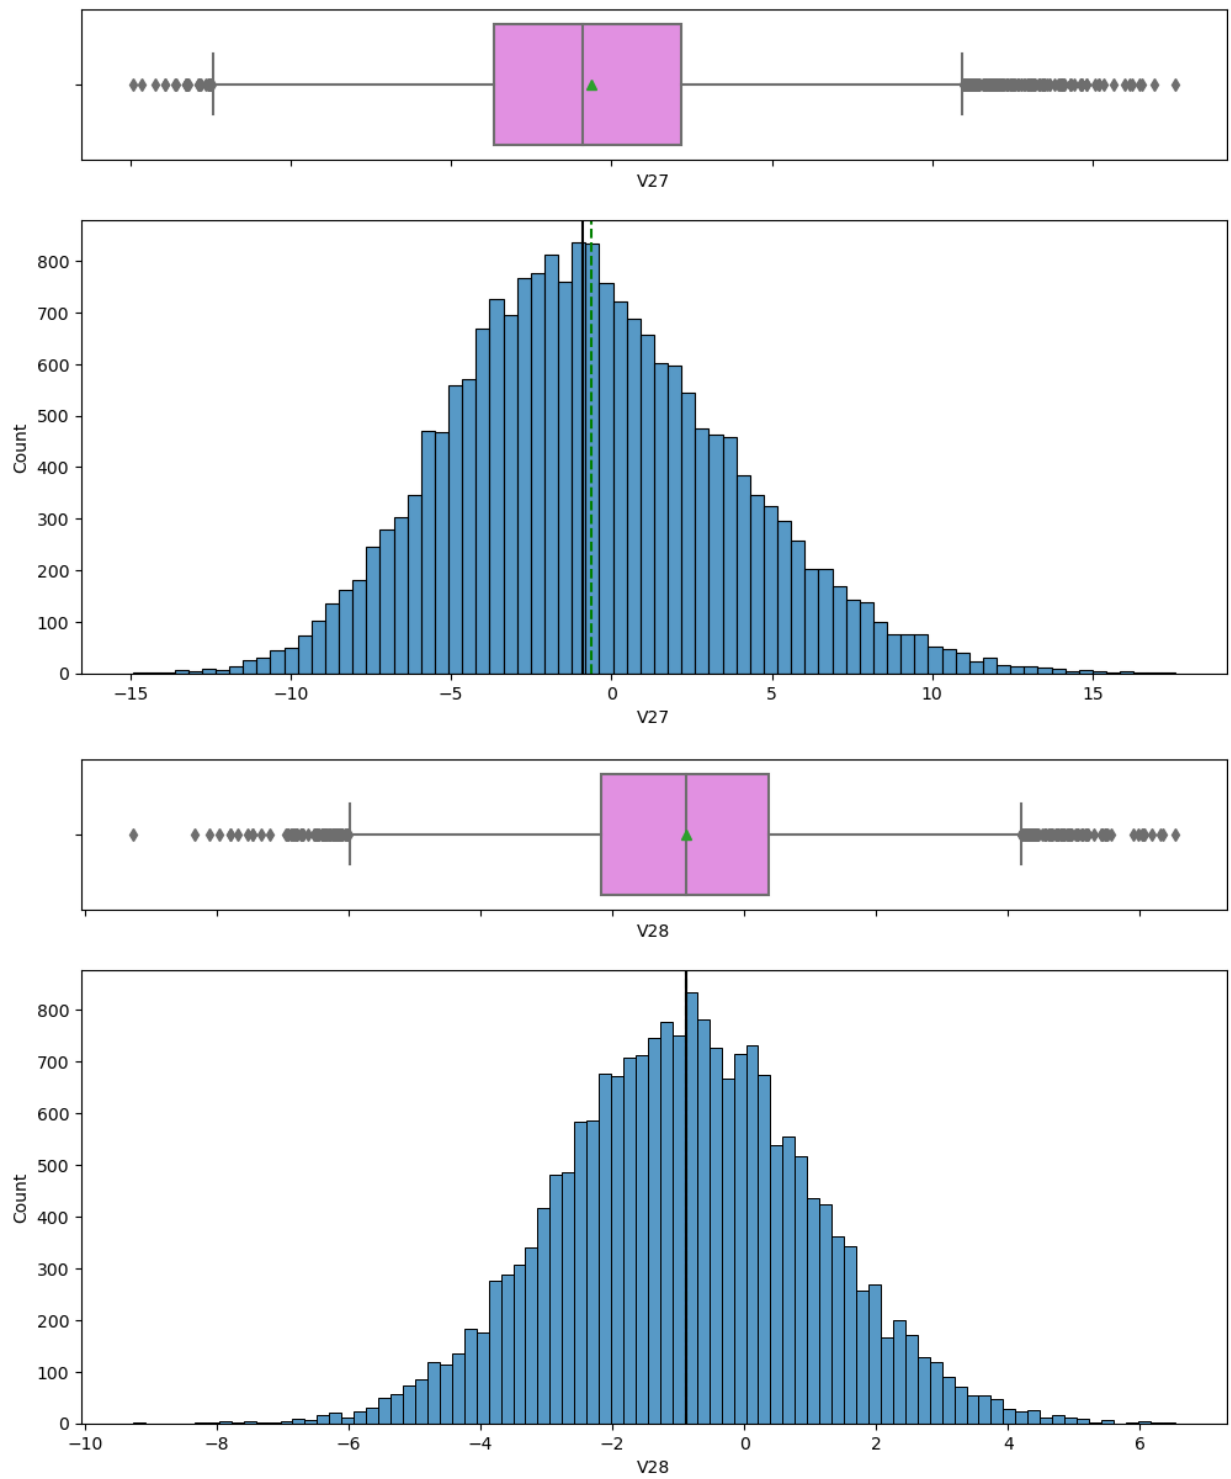


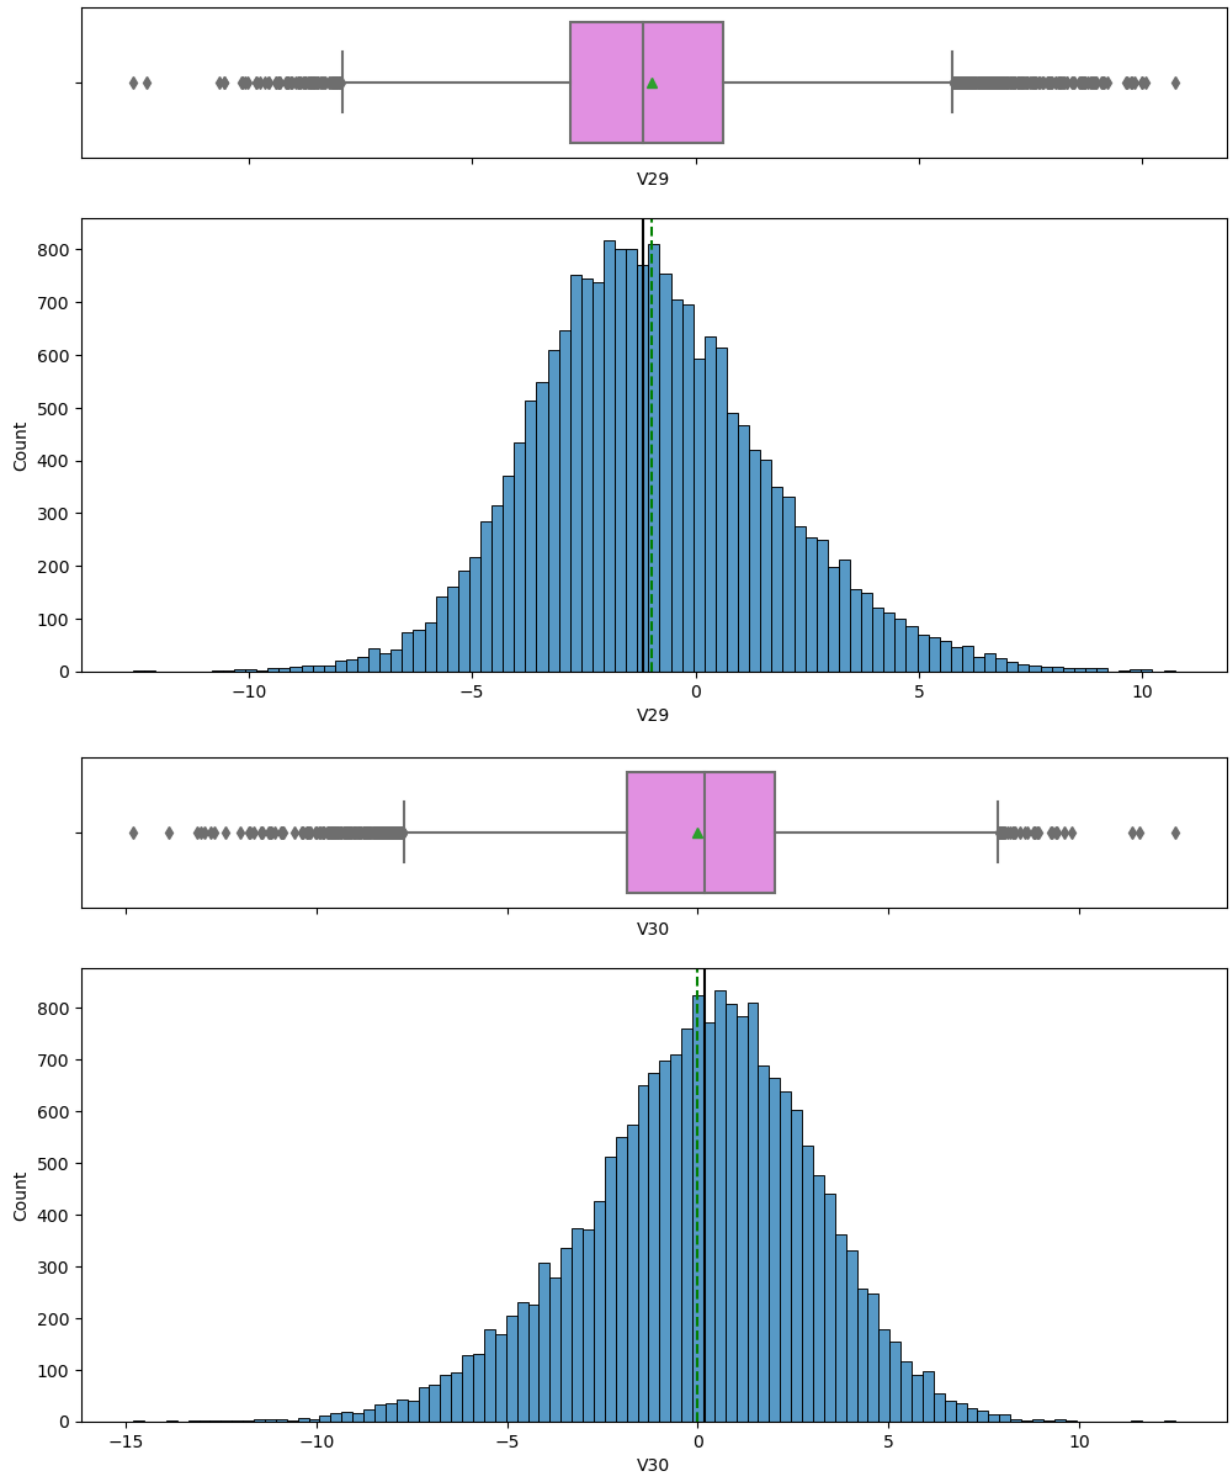


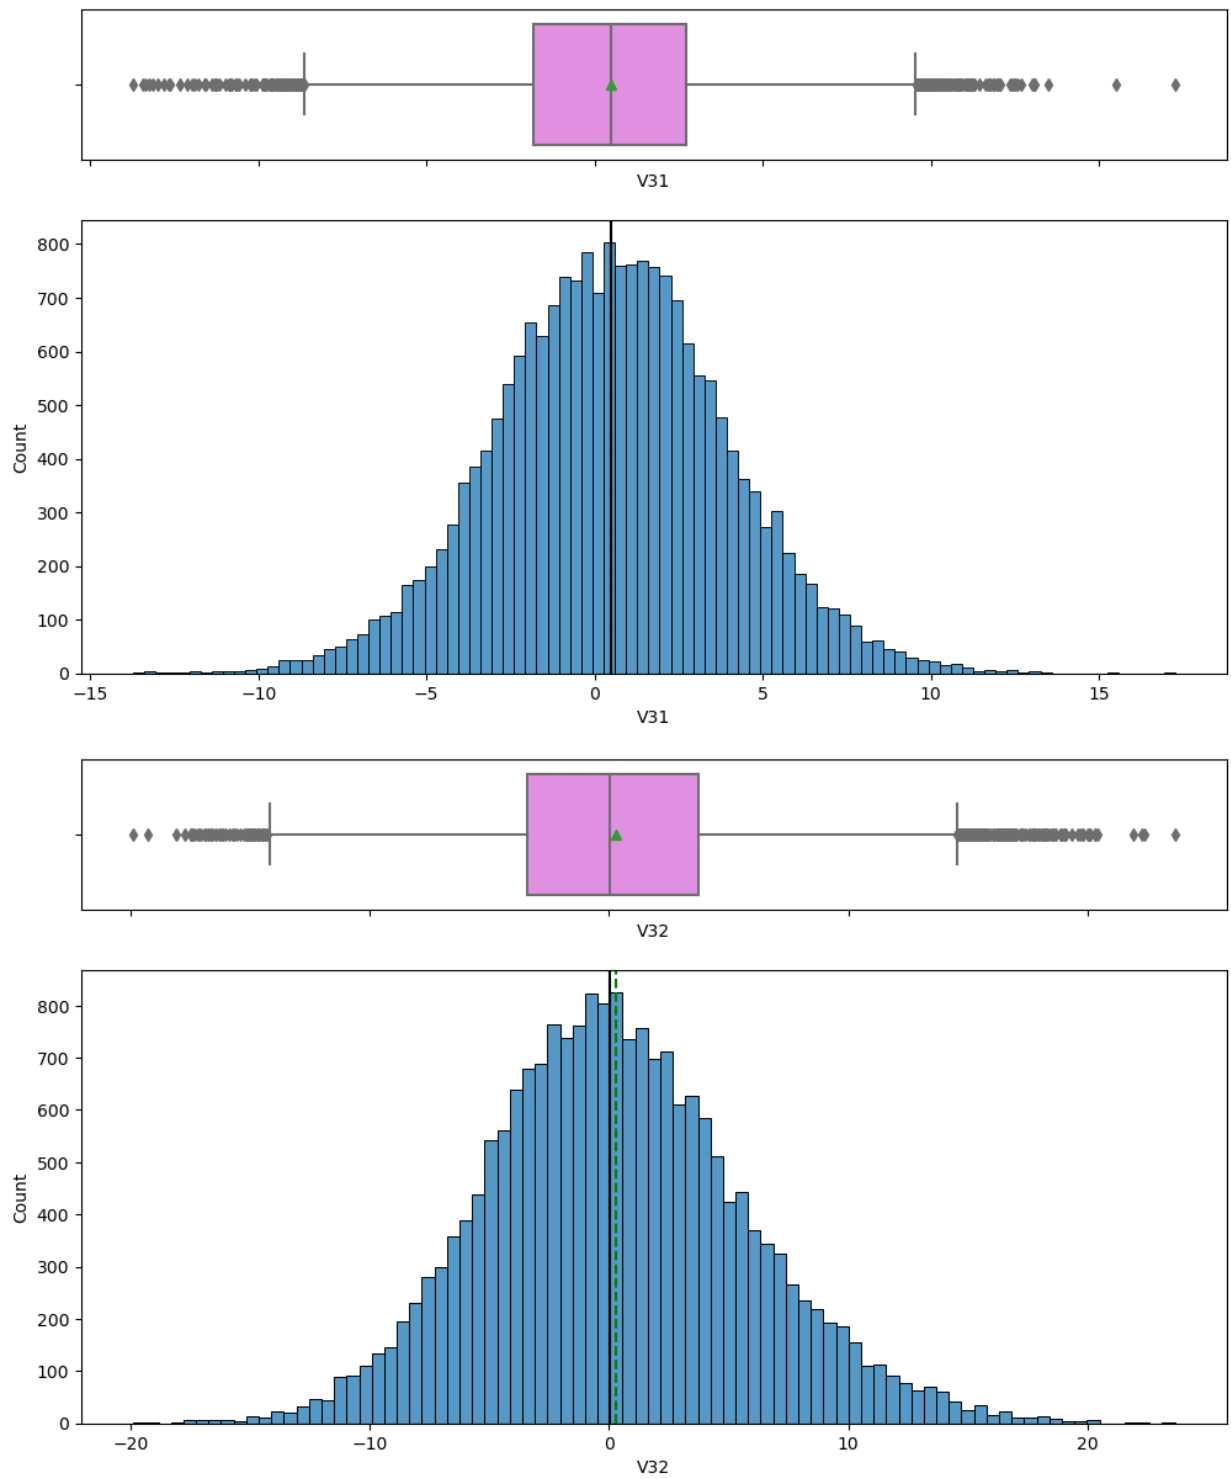


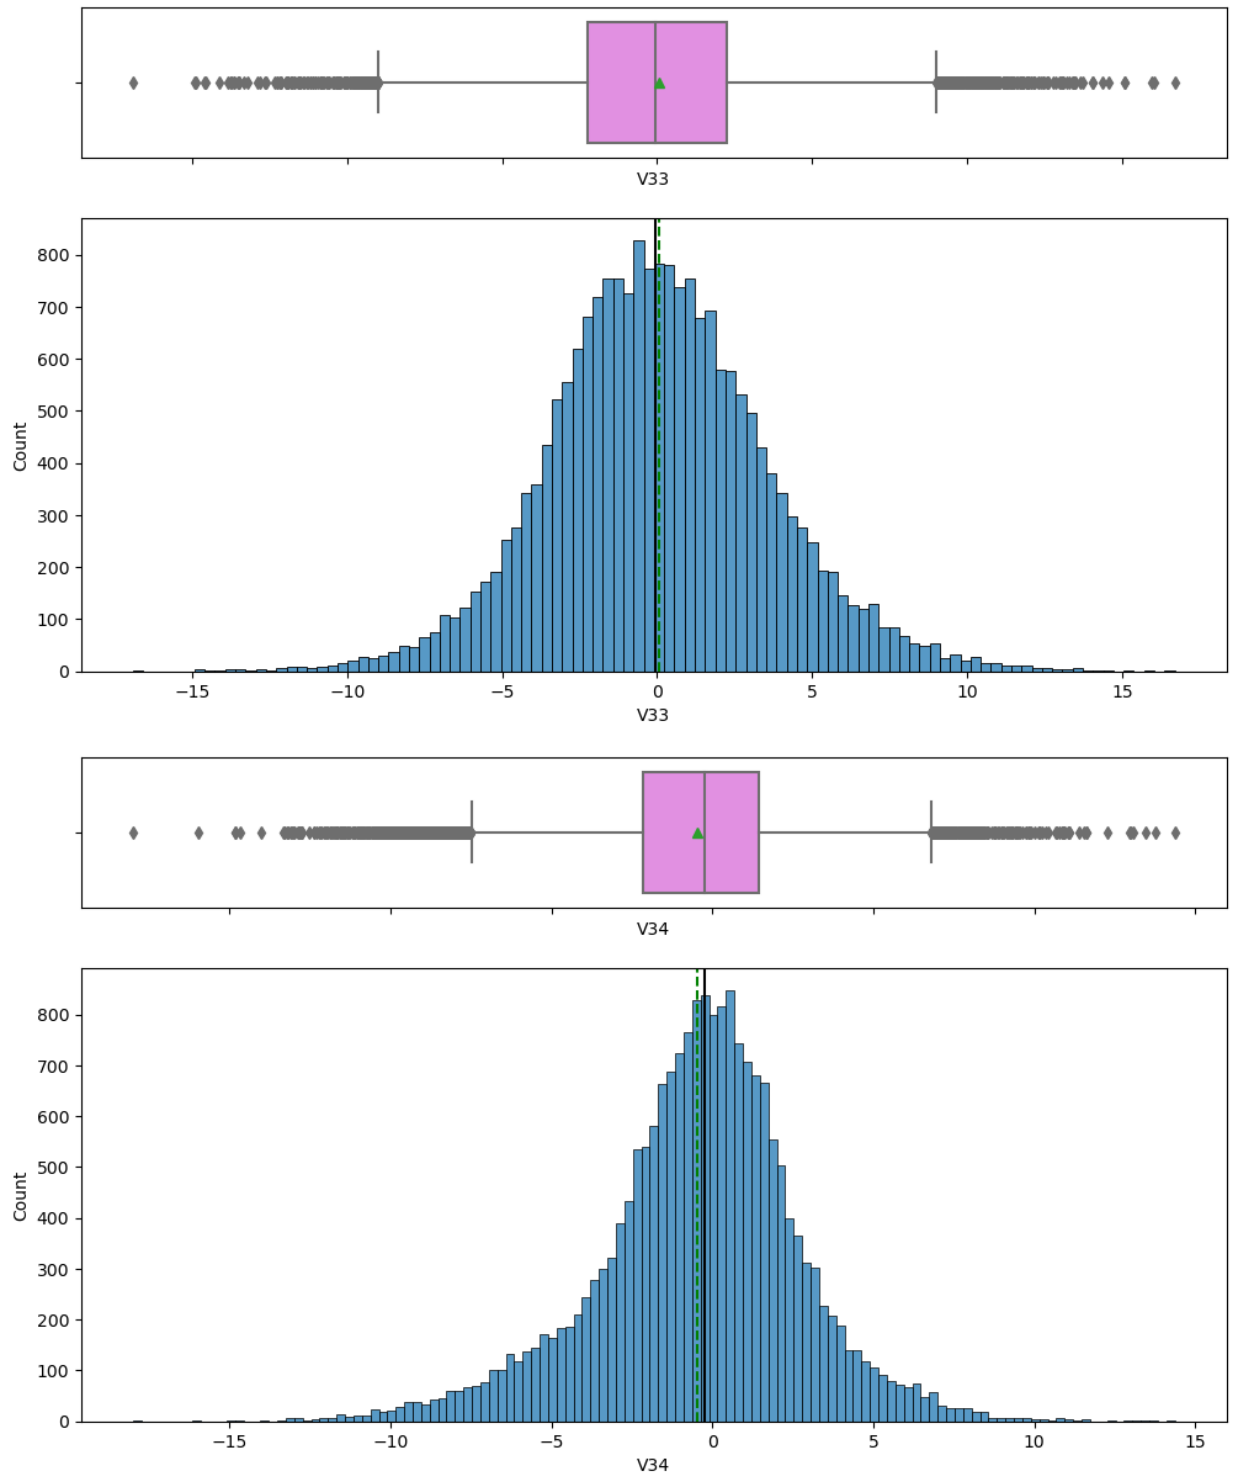


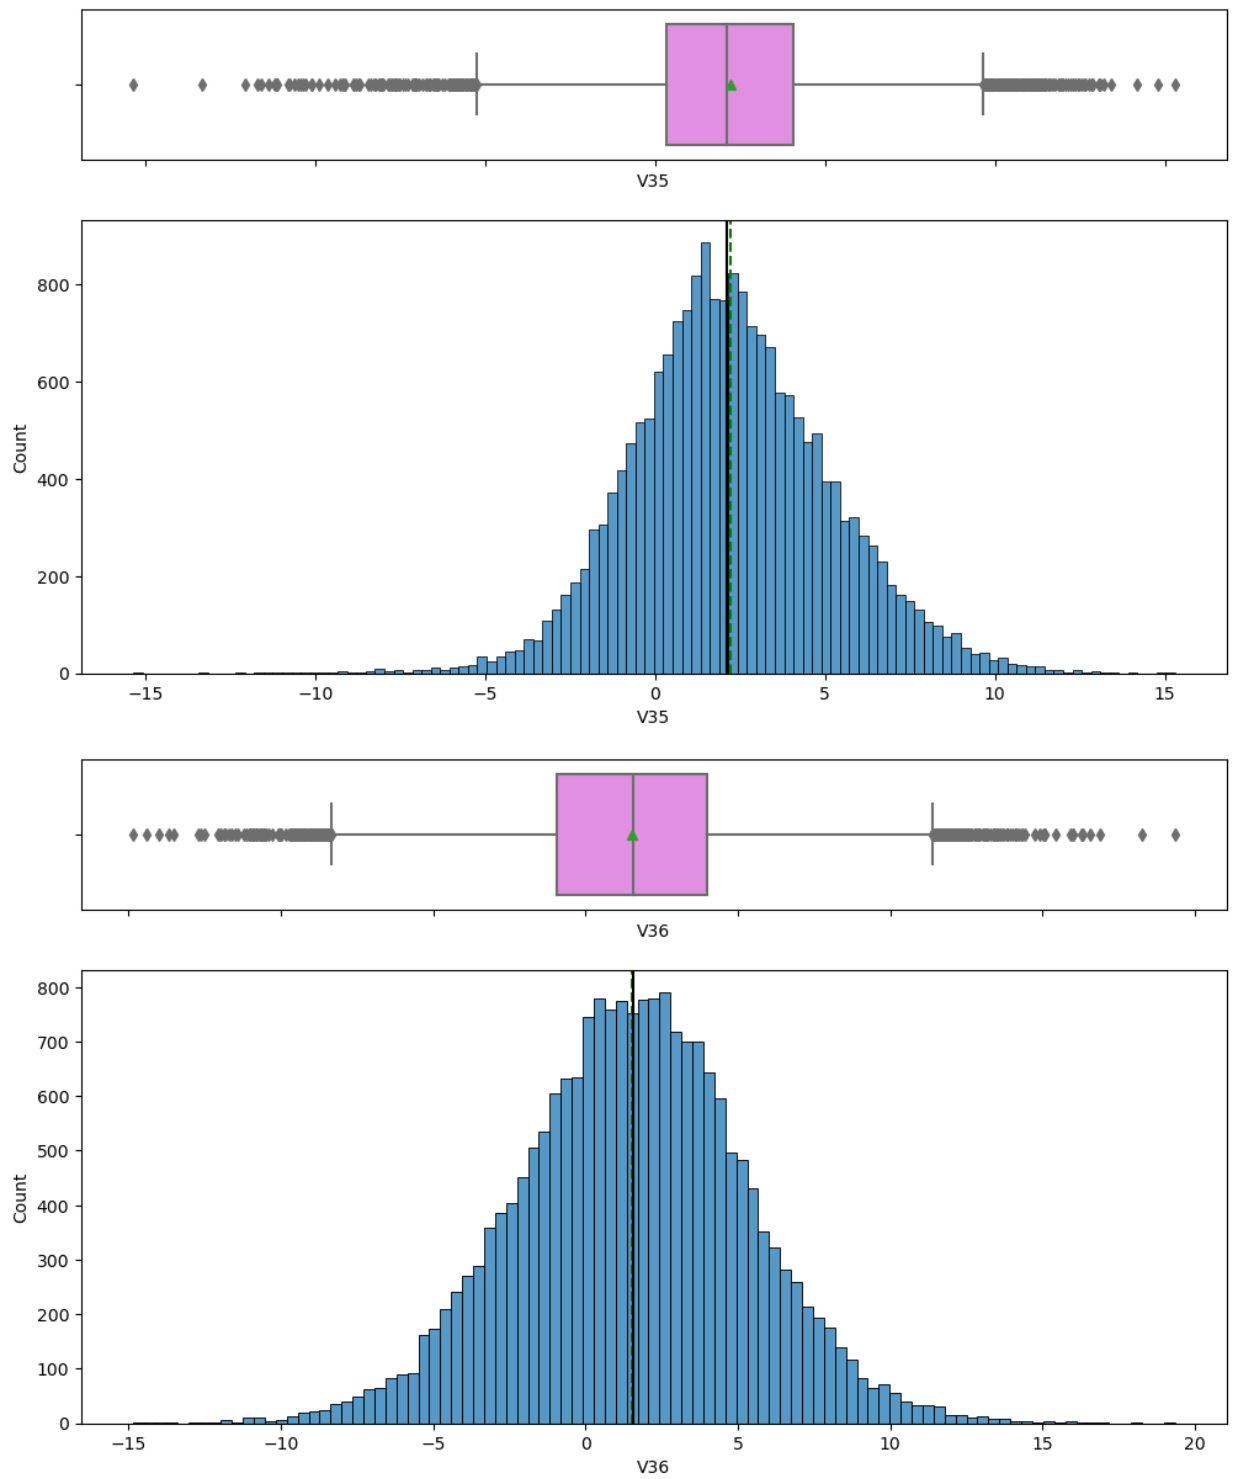


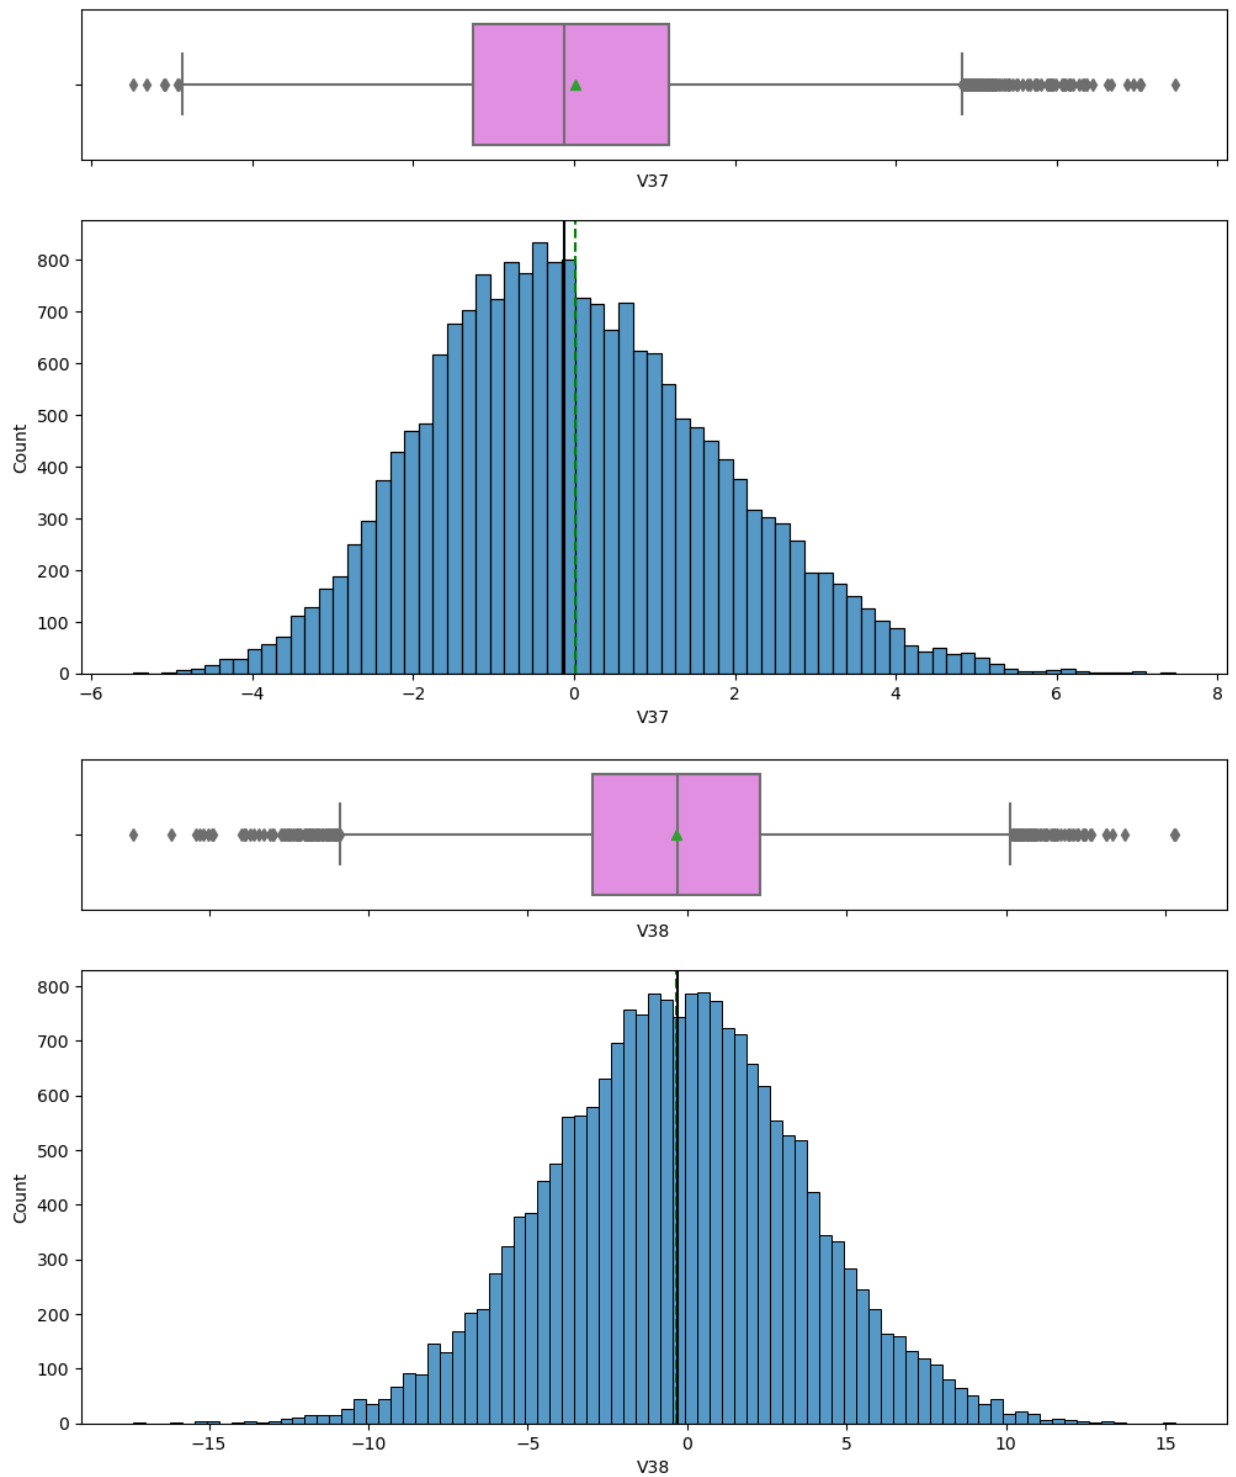


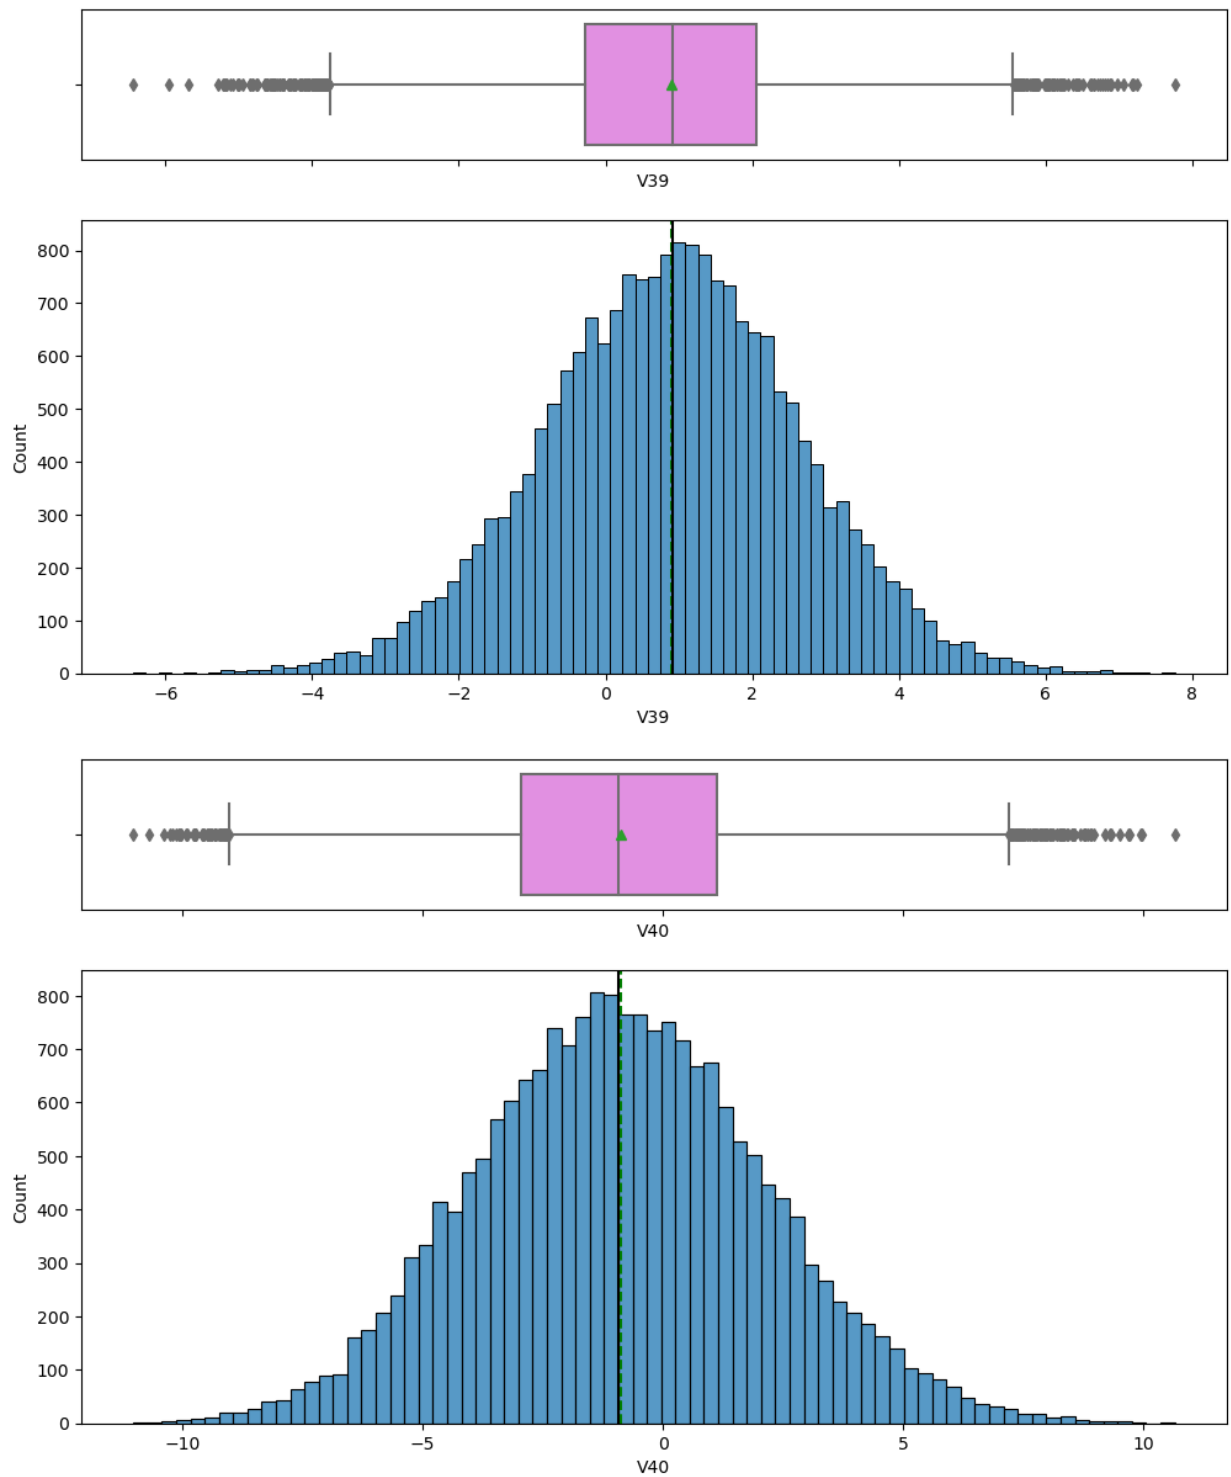















We can see are data is fairly evenly distributed

```
In [16]: df["Target"].value_counts()
```

```
Out[16]: 0    18890
         1     1110
         Name: Target, dtype: int64
```

```
In [17]: df_test["Target"].value_counts()
```

```
Out[17]: 0     4718
         1      282
         Name: Target, dtype: int64
```

Next I will preform prepare the data for model building

```
In [18]: X = df.drop(["Target"], axis=1)
         y = df["Target"]
```

```
In [19]: X_train, X_val, y_train, y_val = train_test_split(
         X, y, test_size=0.25, random_state=1, stratify=y)
```

```
In [20]: X_train.shape
```

```
Out[20]: (15000, 40)
```

```
In [21]: X_val.shape
```

```
Out[21]: (5000, 40)
```

```
In [22]: X_test = df_test.drop(["Target"], axis=1)
         y_test = df_test["Target"]
```

```
In [23]: X_test.shape
```

```
Out[23]: (5000, 40)
```

```
In [24]: imp_mode = SimpleImputer(strategy="median")

# fit and transform the imputer on train data
X_train = pd.DataFrame(imp_mode.fit_transform(X_train), columns=X_train.columns)

# Transform on validation and test data
X_val = pd.DataFrame(imp_mode.fit_transform(X_val), columns=X_train.columns)

# fit and transform the imputer on test data
X_test = pd.DataFrame(imp_mode.fit_transform(X_test), columns=X_train.columns)
```

```
In [25]: print(X_train.isna().sum())
print("-" * 30)
print(X_val.isna().sum())
print("-" * 30)
print(X_test.isna().sum())
```

V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
V29	0
V30	0
V31	0
V32	0
V33	0
V34	0
V35	0
V36	0
V37	0
V38	0
V39	0
V40	0

dtype: int64

V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0

V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
V29	0
V30	0
V31	0
V32	0
V33	0
V34	0
V35	0
V36	0
V37	0
V38	0
V39	0
V40	0

dtype: int64

V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
V29	0
V30	0
V31	0
V32	0
V33	0
V34	0
V35	0
V36	0


```
V37    0
V38    0
V39    0
V40    0
dtype: int64
```

```
In [26]: def model_performance_classification_sklearn(model, predictors, target):
        """
        Function to compute different metrics to check classification model performance

        model: classifier
        predictors: independent variables
        target: dependent variable
        """
        TP= confusion_matrix(target, model.predict(predictors))[1,1]
        FP= confusion_matrix(target, model.predict(predictors))[0,1]
        FN= confusion_matrix(target, model.predict(predictors))[1,0]

        pred = model.predict(predictors)

        acc = accuracy_score(target, pred)
        recall = recall_score(target, pred)
        precision = precision_score(target, pred)
        f1 = f1_score(target, pred)

        df_perf = pd.DataFrame(
            {
                "Accuracy": acc,
                "Recall": recall,
                "Precision": precision,
                "F1": f1,
            },
            index=[0],
        )

        return df_perf
```

```
In [27]: def confusion_matrix_sklearn(model, predictors, target):
        """
        To plot the confusion_matrix with percentages

        model: classifier
        predictors: independent variables
        target: dependent variable
        """
        y_pred = model.predict(predictors)
        cm = confusion_matrix(target, y_pred)
        labels = np.asarray(
            [
                ["{0:0.0f}".format(item) + "\n{0:.2%}".format(item / cm.flatten().sum())]
                for item in cm.flatten()
            ]
        ).reshape(2, 2)

        plt.figure(figsize=(6, 4))
        sns.heatmap(cm, annot=labels, fmt="")
        plt.ylabel("True label")
        plt.xlabel("Predicted label")
```

```
In [28]: scorer = metrics.make_scorer(metrics.recall_score)
```

```
In [29]: models = []

models.append(("Logistic regression", LogisticRegression(random_state=1)))
models.append(("Bagging", BaggingClassifier(random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1)))
models.append(("GBM", GradientBoostingClassifier(random_state=1)))
models.append(("Adaboost", AdaBoostClassifier(random_state=1)))
models.append(("Xgboost", XGBClassifier(random_state=1, eval_metric="logloss")))
models.append(("dtree", DecisionTreeClassifier(random_state=1)))

results = []
names = []

print("\n" "Cross-Validation Performance:" "\n")
for name, model in models:
    scoring = "recall"
    kfold = StratifiedKFold(
        n_splits=5, shuffle=True, random_state=1
    )
    cv_result = cross_val_score(
        estimator=model, X=X_train, y=y_train, scoring=scoring, cv=kfold
    )
    results.append(cv_result)
    names.append(name)
    print("{}: {}".format(name, cv_result.mean() * 100))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train, y_train)
    scores = recall_score(y_val, model.predict(X_val))
    print("{}: {}".format(name, scores))
```

Cross-Validation Performance:

```
Logistic regression: 49.27566553639709
Bagging: 72.1080730106053
Random forest: 72.35192266070268
GBM: 70.66661857008873
Adaboost: 63.09140754635308
Xgboost: 81.00497799581561
dtree: 69.82829521679533
```

Validation Performance:

```
Logistic regression: 0.48201438848920863
Bagging: 0.7302158273381295
Random forest: 0.7266187050359713
GBM: 0.7230215827338129
Adaboost: 0.6762589928057554
Xgboost: 0.8309352517985612
dtree: 0.7050359712230215
```


```
In [30]: fig = plt.figure(figsize=(11,8))

fig.suptitle("Algorithm Comparison")
ax = fig.add_subplot(111)
```

```
plt.boxplot(results)
ax.set_xticklabels(names)

plt.show()
```

Algorithm Comparison



In []: The best performing model now was xgboost

```
In [31]: print("Before UpSampling, counts of label 'Yes': {}".format(sum(y_train == 1)))
print("Before UpSampling, counts of label 'No': {} \n".format(sum(y_train == 0)))

sm = SMOTE(
    sampling_strategy=1, k_neighbors=5, random_state=1
) # Synthetic Minority Over Sampling Technique
X_train_over, y_train_over = sm.fit_resample(X_train, y_train)

print("After UpSampling, counts of label 'Yes': {}".format(sum(y_train_over == 1)))
print("After UpSampling, counts of label 'No': {} \n".format(sum(y_train_over == 0)))

print("After UpSampling, the shape of train_X: {}".format(X_train_over.shape))
print("After UpSampling, the shape of train_y: {} \n".format(y_train_over.shape))
```

Before UpSampling, counts of label 'Yes': 832
 Before UpSampling, counts of label 'No': 14168

After UpSampling, counts of label 'Yes': 14168
 After UpSampling, counts of label 'No': 14168

After UpSampling, the shape of train_X: (28336, 40)
 After UpSampling, the shape of train_y: (28336,)

Next I will build my overfitting models

```
In [32]: models = []

models.append(("Logistic regression", LogisticRegression(random_state=1)))
models.append(("Bagging", BaggingClassifier(random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1)))
models.append(("GBM", GradientBoostingClassifier(random_state=1)))
models.append(("Adaboost", AdaBoostClassifier(random_state=1)))
models.append(("Xgboost", XGBClassifier(random_state=1, eval_metric="logloss")))
models.append(("dtree", DecisionTreeClassifier(random_state=1)))

results = []
names = []

print("\n" "Cross-Validation Performance:" "\n")
for name, model in models:
    scoring = "recall"
    kfold = StratifiedKFold(
        n_splits=5, shuffle=True, random_state=1
    )
    cv_result = cross_val_score(
        estimator=model, X=X_train_over, y=y_train_over, scoring=scoring, cv=kfold
    )
    results.append(cv_result)
    names.append(name)
    print("{}: {}".format(name, cv_result.mean() * 100))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train_over, y_train_over)
    scores = recall_score(y_val, model.predict(X_val))
    print("{}: {}".format(name, scores))
```

Cross-Validation Performance:


Logistic regression: 88.3963699328486
Bagging: 97.62141471581656
Random forest: 98.39075260047615
GBM: 92.56068151319724
Adaboost: 89.78689011775472
Xgboost: 98.91305241357217
dtree: 97.20494245534968

Validation Performance:

Logistic regression: 0.8489208633093526
Bagging: 0.8345323741007195
Random forest: 0.8489208633093526
GBM: 0.8776978417266187
Adaboost: 0.8561151079136691
Xgboost: 0.8669064748201439
dtree: 0.7769784172661871

```
In [33]: fig = plt.figure(figsize=(11,8))  
  
fig.suptitle("Algorithm Comparison")  
ax = fig.add_subplot(111)  
  
plt.boxplot(results)  
ax.set_xticklabels(names)  
  
plt.show()
```

Algorithm Comparison



My best performing models now were adaboost and GBM

```
In [34]: rus = RandomUnderSampler(random_state=1)
X_train_un, y_train_un = rus.fit_resample(X_train, y_train)
```

```
In [35]: print("Before Under Sampling, counts of label 'Yes': {}".format(sum(y_train == 1)))
print("Before Under Sampling, counts of label 'No': {} \n".format(sum(y_train == 0)))

print("After Under Sampling, counts of label 'Yes': {}".format(sum(y_train_un == 1)))
print("After Under Sampling, counts of label 'No': {} \n".format(sum(y_train_un == 0)))

print("After Under Sampling, the shape of train_X: {}".format(X_train_un.shape))
print("After Under Sampling, the shape of train_y: {} \n".format(y_train_un.shape))
```

```
Before Under Sampling, counts of label 'Yes': 832
Before Under Sampling, counts of label 'No': 14168
```

```
After Under Sampling, counts of label 'Yes': 832
After Under Sampling, counts of label 'No': 832
```

```
After Under Sampling, the shape of train_X: (1664, 40)
After Under Sampling, the shape of train_y: (1664,)
```

Next I will build my underfitting models

```
In [36]: models = []

models.append(("Logistic regression", LogisticRegression(random_state=1)))
models.append(("Bagging", BaggingClassifier(random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1)))
models.append(("GBM", GradientBoostingClassifier(random_state=1)))
models.append(("Adaboost", AdaBoostClassifier(random_state=1)))
models.append(("Xgboost", XGBClassifier(random_state=1, eval_metric="logloss")))
models.append(("dtree", DecisionTreeClassifier(random_state=1)))

results = []
names = []

print("\n" "Cross-Validation Performance:" "\n")
for name, model in models:
    scoring = "recall"
    kfold = StratifiedKFold(
        n_splits=5, shuffle=True, random_state=1
    )
    cv_result = cross_val_score(
        estimator=model, X=X_train_un, y=y_train_un, scoring=scoring, cv=kfold
    )
    results.append(cv_result)
    names.append(name)
    print("{}: {}".format(name, cv_result.mean() * 100))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train_un, y_train_un)
    scores = recall_score(y_val, model.predict(X_val))
    print("{}: {}".format(name, scores))
```

Cross-Validation Performance:

```
Logistic regression: 87.26138085275232
Bagging: 86.41945025611427
Random forest: 90.38669648654498
GBM: 89.78572974532861
Adaboost: 86.6611355602049
Xgboost: 90.14717552846115
dtree: 86.17776495202367
```

Validation Performance:

```
Logistic regression: 0.8525179856115108
Bagging: 0.8705035971223022
Random forest: 0.8920863309352518
GBM: 0.8884892086330936
Adaboost: 0.8489208633093526
Xgboost: 0.89568345323741
dtree: 0.841726618705036
```

```
In [37]: fig = plt.figure(figsize=(11,8))


fig.suptitle("Algorithm Comparison")
ax = fig.add_subplot(111)

plt.boxplot(results)
```

```
ax.set_xticklabels(names)

plt.show()
```

Algorithm Comparison



My best performing model was random forest

I will now tune my models

```
In [ ]: %%time

#defining model
model = AdaBoostClassifier(random_state=1)

#Parameter grid to pass in RandomizedSearchCV
param_grid = {'n_estimators':[100,150,200],
              'learning_rate':[0.2,0.05],
              'base_estimator':[DecisionTreeClassifier(max_depth=1, random_state=1), Dec

}

#Calling RandomizedSearchCV
Randomized_cv = RandomizedSearchCV(estimator=model, param_distributions=param_grid, sc

#Fitting parameters in RandomizedSearchCV
Randomized_cv.fit(X_train_over,y_train_over)
```



```
print("Best parameters are {} with CV score={:}" .format(Randomized_cv.best_params_,Ra
```

```
In [39]: tune_ada = AdaBoostClassifier(n_estimators= 200, learning_rate= 0.2, base_estimator= D
tune_ada.fit(X_train_over,y_train_over)
```

```
Out[39]: ▸ AdaBoostClassifier
▸ base_estimator: DecisionTreeClassifier
▸ DecisionTreeClassifier
```

```
In [40]: ada_train_perf = model_performance_classification_sklearn(tune_ada,X_train_over,y_train_over)
ada_train_perf
```

```
Out[40]:
```

	Accuracy	Recall	Precision	F1
0	0.992	0.988	0.995	0.992

```
In [41]: ada_val_perf = model_performance_classification_sklearn(tune_ada,X_val,y_val)
ada_val_perf
```

```
Out[41]:
```

	Accuracy	Recall	Precision	F1
0	0.979	0.849	0.789	0.818

```
In [42]: model = RandomForestClassifier(random_state=1)

# Parameter grid to pass in RandomizedSearchCV
param_grid = {'n_estimators':[200,250,300],
              'min_samples_leaf': np.arange(1,4),
              'max_features': [np.arange(0.3,0.6,0.1), 'sqrt'],
              'max_samples': np.arange(0.4,0.7,0.1)
              }

# Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_grid,
    n_iter=50,
    n_jobs=-1,
    scoring=scorer,
    cv=5,
    random_state=1,
)

# Fitting parameters in RandomizedSearchCV
randomized_cv.fit(X_train_un, y_train_un)

print(
    "Best parameters are {} with CV score={:}" .format(
        randomized_cv.best_params_, randomized_cv.best_score_
    )
)
```

Best parameters are {'n_estimators': 300, 'min_samples_leaf': 2, 'max_samples': 0.5, 'max_features': 'sqrt'} with CV score=0.8990116153235697:

```
In [43]: tune_rf = RandomForestClassifier(n_estimators= 300, min_samples_leaf= 2, max_samples=
tune_rf.fit(X_train_un,y_train_un)
```

```
Out[43]: ▼ RandomForestClassifier
RandomForestClassifier(max_samples=0.5, min_samples_leaf=2, n_estimators=300)
```

```
In [44]: rf_train_perf = model_performance_classification_sklearn(tune_rf,X_train_un,y_train_un)
rf_train_perf
```

```
Out[44]:
```

	Accuracy	Recall	Precision	F1
0	0.962	0.934	0.990	0.961

```
In [45]: rf_val_perf = model_performance_classification_sklearn(tune_rf,X_val,y_val)
rf_val_perf
```

```
Out[45]:
```

	Accuracy	Recall	Precision	F1
0	0.935	0.881	0.457	0.602

```
In [46]: model = GradientBoostingClassifier(random_state=1)

# Parameter grid to pass in RandomizedSearchCV
param_grid = {'n_estimators':[200,250,300],
              'learning_rate': np.arange(.05,.2,1),
              'max_features': np.arange(0.5,0.7),
              'subsample':np.arange(0.5,0.7)
              }

# Calling RandomizedSearchCV
randomized_cv = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_grid,
    n_iter=50,
    n_jobs=-1,
    scoring=scorer,
    cv=5,
    random_state=1,
)

# Fitting parameters in RandomizedSearchCV
randomized_cv.fit(X_train_over, y_train_over)

print(
    "Best parameters are {} with CV score={}".format(
        randomized_cv.best_params_, randomized_cv.best_score_
    )
)
```

Best parameters are {'subsample': 0.5, 'n_estimators': 300, 'max_features': 0.5, 'learning_rate': 0.05} with CV score=0.9335123572593496:

```
In [47]: tune_gb = GradientBoostingClassifier(n_estimators= 300, subsample= 0.5, learning_rate=
tune_gb.fit(X_train_over, y_train_over)
```

```
Out[47]: GradientBoostingClassifier
GradientBoostingClassifier(learning_rate=0.05, max_features=0.5,
n_estimators=300, subsample=0.5)
```

```
In [48]: gb_train_perf = model_performance_classification_sklearn(tune_gb,X_train_over,y_train_
gb_train_perf
```

```
Out[48]:
```

	Accuracy	Recall	Precision	F1
0	0.960	0.937	0.981	0.959

```
In [49]: gb_val_perf = model_performance_classification_sklearn(tune_gb,X_val,y_val)
gb_val_perf
```

```
Out[49]:
```

	Accuracy	Recall	Precision	F1
0	0.971	0.878	0.683	0.769

```
In [50]: models_train_comp_df = pd.concat(
[
    gb_val_perf.T,
    rf_val_perf.T,
    ada_val_perf.T,
],
axis=1,
)
models_train_comp_df.columns = [
    "Gradient Boosting",
    "Random Forest",
    "Ada",
]
print("Training performance comparison:")
models_train_comp_df
```

Training performance comparison:

```
Out[50]:
```


	Gradient Boosting	Random Forest	Ada
Accuracy	0.971	0.935	0.979
Recall	0.878	0.881	0.849
Precision	0.683	0.457	0.789
F1	0.769	0.602	0.818

My best performing model was Ada so thats what I'll use to build my final model

```
In [51]: feature_names = X_train.columns
importances = tune_ada.feature_importances_
indices = np.argsort(importances)

plt.figure(figsize=(12, 12))
```

```
plt.title("Feature Importances")
plt.barh(range(len(indices)), importances[indices], color="violet", align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()
```

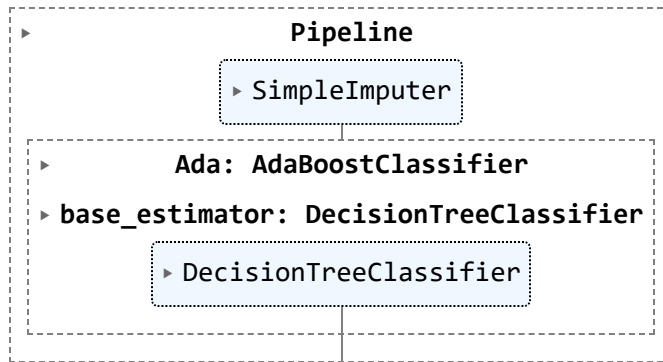


The 3 most important variables are v30, v9, and v18.

```
In [54]: # Creating new pipeline with best parameters
Final_pipe = Pipeline(
    steps=[
        ("imputer", SimpleImputer(strategy='median')),
        (
            "Ada",
            AdaBoostClassifier(
                n_estimators= 200, learning_rate= 0.2, base_estimator= DecisionTreeClas
            )
        )
    ])
])
```

```
# Fit the model on training data
Final_pipe.fit(X_train, y_train)
```

Out[54]:



Next I will prepare my final model

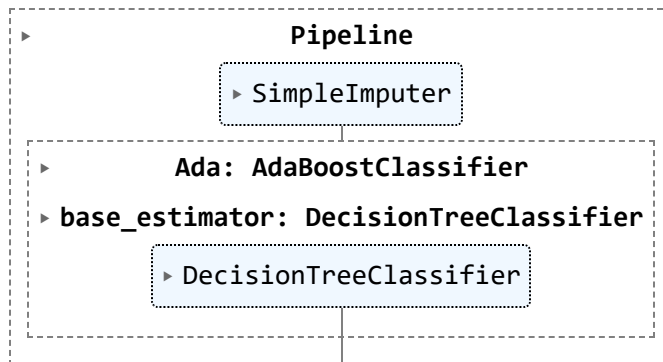
```
In [60]: X1 = data.drop(["Target"], axis=1)
y1 = data["Target"]
Xt1 = data_test.drop(["Target"], axis=1)
yt1 = data_test["Target"]
```

```
In [56]: imputer = SimpleImputer(strategy='median')
X1 = imputer.fit_transform(X1)
```

```
In [57]: SM = SMOTE(sampling_strategy=1, k_neighbors=5, random_state=1)
X_overf, y_overf = sm.fit_resample(X1, y1)
```

```
In [58]: Final_pipe.fit(X_overf, y_overf)
```

Out[58]:



```
In [59]: Final_pipe_perf = model_performance_classification_sklern(Final_pipe, X_test, y_test)
Final_pipe_perf
```

Out[59]:

	Accuracy	Recall	Precision	F1
0	0.978	0.851	0.774	0.811

Ada using oversampling will yeild the best model

This model will better help predict wind turbine failures

v30, v9, and v18 are the most important variables and play a key role in this model